

Elements Of Data Science - F2021

Week 12: Time Series, Imbalanced Classes, Data Processing and Delivery

12/6/2021

TODOs

- Readings:
 - Recommended: DSFS: Chap 23: Databases and SQL
 - Final Review Sheet
- HW4, **Due Saturday December 11th 11:59pm ET**
- Quiz 11, **Due Sunday December 12th, 11:59pm ET**
- Final
 - Released Monday night December 13th, 11:59pm
 - **Due Wednesday December 15th, 11:59pm ET**
 - Have 24hrs after starting exam to finish
 - 30-40 questions (fill in the blank/multiple choice/short answer)
 - Online via Gradescope
 - Questions asked/answered **privately** via Ed
 - Open-book, open-note, open-python

Today

- Finish Time Series
- Imbalanced Classes
- Data processing and delivery

Questions?

Environment Setup

Environment Setup

```
In [1]: import numpy
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('darkgrid')
%matplotlib inline
```

Time Series in Python so far:

- `datetime .date .time .datetime .timedelta`
- format with `.strftime()`
- parse time with `pd.to_datetime()`
- `pandas Timestamp Timedelta DatetimeIndex`
- Indexing with `DatetimeIndex`
- Frequencies
- Timezones

Next: Operations on Time Series data

- Shifting
- Resampling
- Moving Windows

Shifting

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

Shifting

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [2]: ts = pd.Series([1,2,8],  
                        index=pd.date_range('1/1/2019', periods=3, freq='M'))  
ts
```

```
Out[2]: 2019-01-31    1  
        2019-02-28    2  
        2019-03-31    8  
        Freq: M, dtype: int64
```

Shifting

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [2]: ts = pd.Series([1,2,8],  
                        index=pd.date_range('1/1/2019', periods=3, freq='M'))  
ts
```

```
Out[2]: 2019-01-31    1  
        2019-02-28    2  
        2019-03-31    8  
        Freq: M, dtype: int64
```

```
In [3]: ts.shift(1) # last month's value
```

```
Out[3]: 2019-01-31    NaN  
        2019-02-28    1.0  
        2019-03-31    2.0  
        Freq: M, dtype: float64
```

Shifting

- percent change :
 - $(\text{new_value} - \text{old_value}) / \text{old_value}$
 - $(\text{new_value} / \text{old_value}) - 1$

Shifting

- percent change :
 - $(\text{new_value} - \text{old_value}) / \text{old_value}$
 - $(\text{new_value} / \text{old_value}) - 1$

```
In [4]: # multiply by 100 to turn into a percent
        ((ts / ts.shift(1)) - 1) * 100
```

```
Out[4]: 2019-01-31      NaN
        2019-02-28    100.0
        2019-03-31    300.0
        Freq: M, dtype: float64
```

Example Dataset: Twitter Stock

Example Dataset: Twitter Stock

```
In [5]: #from pandas_datareader import data
#df_twtr = data.DataReader('TWTR', start='2015', end='12/3/2021', data_source='yahoo')
#df_twtr.to_csv('../data/twtr_20150102-20211203.csv')
df_twtr = pd.read_csv('../data/twtr_20150102-20211203.csv', parse_dates=['Date'], index_col='Date')
df_twtr.head(3)
```

Out[5]:

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.740002	35.540001	36.230000	36.560001	12062500	36.560001
2015-01-05	37.110001	35.639999	36.259998	36.380001	15062700	36.380001
2015-01-06	39.450001	36.040001	36.270000	38.759998	33050800	38.759998

Example Dataset: Twitter Stock

```
In [5]: #from pandas_datareader import data
#df_twtr = data.DataReader('TWTR', start='2015', end='12/3/2021', data_source='yahoo')
#df_twtr.to_csv('../data/twtr_20150102-20211203.csv')
df_twtr = pd.read_csv('../data/twtr_20150102-20211203.csv', parse_dates=['Date'], index_col='Date')
df_twtr.head(3)
```

Out[5]:

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.740002	35.540001	36.230000	36.560001	12062500	36.560001
2015-01-05	37.110001	35.639999	36.259998	36.380001	15062700	36.380001
2015-01-06	39.450001	36.040001	36.270000	38.759998	33050800	38.759998

```
In [6]: df_twtr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1744 entries, 2015-01-02 to 2021-12-03
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   High        1744 non-null   float64
1   Low         1744 non-null   float64
2   Open        1744 non-null   float64
3   Close       1744 non-null   float64
4   Volume      1744 non-null   int64
5   Adj Close   1744 non-null   float64
dtypes: float64(5), int64(1)
memory usage: 95.4 KB
```

Example Dataset: Twitter Stock

Example Dataset: Twitter Stock

```
In [7]: fig, ax = plt.subplots(1, 1, figsize=(24, 10))
df_twtr[['Close']].plot(ax=ax, marker='x');
ax.set_ylabel('stock price');
```



Shifting Example: Percent Change Twitter Close

Shifting Example: Percent Change Twitter Close

```
In [8]: ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3) # # (today / yesterday) - 1
```

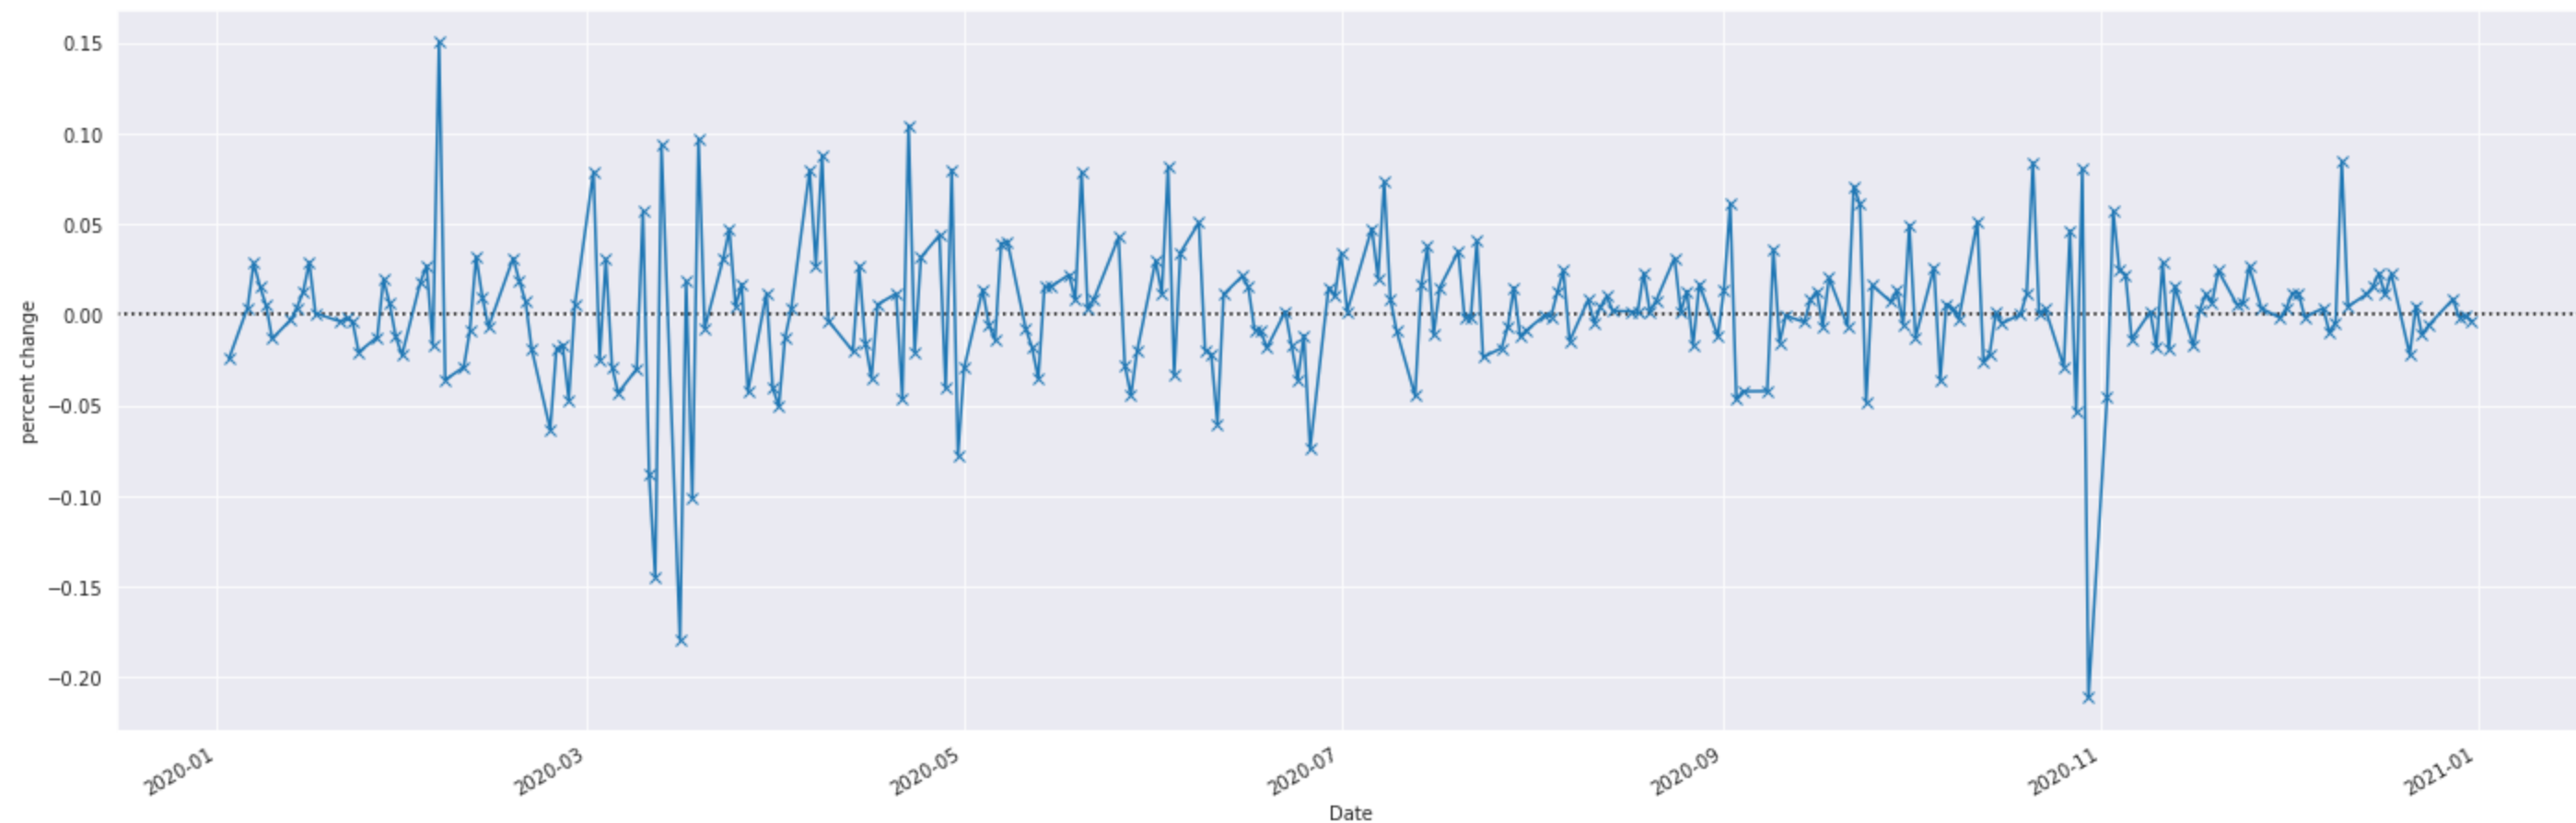
```
Out[8]: Date  
2021-12-01    -0.025489  
2021-12-02    -0.003970  
2021-12-03    -0.013599  
Name: Close, dtype: float64
```

Shifting Example: Percent Change Twitter Close

```
In [8]: ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3) # # (today / yesterday) - 1
```

```
Out[8]: Date
2021-12-01    -0.025489
2021-12-02    -0.003970
2021-12-03    -0.013599
Name: Close, dtype: float64
```

```
In [9]: # plot percent change of close in 2020
fig,ax = plt.subplots(1,1,figsize=(24,8))
close_2020 = df_twtr.loc['2020', 'Close']
((close_2020 / close_2020.shift(1)) - 1).plot(marker='x',ax=ax,zorder=2);
ax.axhline(ls=':',c='k',zorder=1)
ax.set_ylabel('percent change');
```



Resampling

- Convert from one frequency to another
- **Downsampling**
 - from higher to lower (day to month)
 - need to aggregate
- **Upsampling**
 - from lower to higher (month to day)
 - need to fill missing
- **Can also be used to set frequency from None**

Resampling: Initialize Frequency

Resampling: Initialize Frequency

```
In [10]: df_twtr.index
```

```
Out[10]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...,  
                        '2021-11-19', '2021-11-22', '2021-11-23', '2021-11-24',  
                        '2021-11-26', '2021-11-29', '2021-11-30', '2021-12-01',  
                        '2021-12-02', '2021-12-03'],  
                        dtype='datetime64[ns]', name='Date', length=1744, freq=None)
```

Resampling: Initialize Frequency

```
In [10]: df_twtr.index
```

```
Out[10]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...  
                        '2021-11-19', '2021-11-22', '2021-11-23', '2021-11-24',  
                        '2021-11-26', '2021-11-29', '2021-11-30', '2021-12-01',  
                        '2021-12-02', '2021-12-03'],  
                      dtype='datetime64[ns]', name='Date', length=1744, freq=None)
```

```
In [11]: df_twtr_B = df_twtr.resample('B').asfreq() # set frequency to business day  
df_twtr_B.index
```

```
Out[11]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...  
                        '2021-11-22', '2021-11-23', '2021-11-24', '2021-11-25',  
                        '2021-11-26', '2021-11-29', '2021-11-30', '2021-12-01',  
                        '2021-12-02', '2021-12-03'],  
                      dtype='datetime64[ns]', name='Date', length=1806, freq='B')
```


Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [12]: df_twtr_BQ = df_twtr_B.resample('BQ')  
df_twtr_BQ
```

```
Out[12]: <pandas.core.resample.DatetimeIndexResampler object at 0x7f88c261bb80>
```

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [12]: df_twtr_BQ = df_twtr_B.resample('BQ')  
df_twtr_BQ
```

```
Out[12]: <pandas.core.resample.DatetimeIndexResampler object at 0x7f88c261bb80>
```

```
In [13]: str(df_twtr_BQ)
```

```
Out[13]: 'DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]'
```

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [12]: df_twtr_BQ = df_twtr_B.resample('BQ')
df_twtr_BQ
```

```
Out[12]: <pandas.core.resample.DatetimeIndexResampler object at 0x7f88c261bb80>
```

```
In [13]: str(df_twtr_BQ)
```

```
Out[13]: 'DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]'
```

```
In [14]: df_twtr_BQ.mean().head(3)
```

```
Out[14]:
```

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-03-31	45.080328	43.552459	44.228688	44.335574	2.084619e+07	44.335574
2015-06-30	41.634921	40.385079	41.173492	40.874603	2.232030e+07	40.874603
2015-09-30	30.638281	29.420625	30.047812	30.000625	2.031210e+07	30.000625

Resampling: Downsampling

Resampling: Downsampling

```
In [15]: fig,ax = plt.subplots(1,1,figsize=(24,8))
df_twtr_B.Close.plot(style='-', label='by B',ax=ax)
df_twtr_BQ.Close.mean().plot(style='--',marker='x',label='by BQ',ax=ax)
plt.legend(loc='upper right');
```



Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [16]: df_twtr_B.index[:3]
```

```
Out[16]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```


Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [16]: df_twtr_B.index[:3]
```

```
Out[16]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [17]: df_twtr_B.Close.resample('H').asfreq().iloc[0:3]
```

```
Out[17]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00         NaN
2015-01-02 02:00:00         NaN
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [16]: df_twtr_B.index[:3]
```

```
Out[16]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [17]: df_twtr_B.Close.resample('H').asfreq().iloc[0:3]
```

```
Out[17]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00         NaN
2015-01-02 02:00:00         NaN
Freq: H, Name: Close, dtype: float64
```

```
In [18]: df_twtr_B.Close.resample('H').asfreq().iloc[70:73]
```

```
Out[18]: Date
2015-01-04 22:00:00         NaN
2015-01-04 23:00:00         NaN
2015-01-05 00:00:00    36.380001
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- `ffill()` : Forward Fill

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [19]: df_twtr_B.Close.resample('H').ffill().head(3)
```

```
Out[19]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00    36.560001
2015-01-02 02:00:00    36.560001
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [19]: df_twtr_B.Close.resample('H').ffill().head(3)
```

```
Out[19]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00    36.560001
2015-01-02 02:00:00    36.560001
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [19]: df_twtr_B.Close.resample('H').ffill().head(3)
```

```
Out[19]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00    36.560001
2015-01-02 02:00:00    36.560001
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill

```
In [20]: df_twtr_B.Close.resample('H').bfill().head(3)
```

```
Out[20]: Date
2015-01-02 00:00:00    36.560001
2015-01-02 01:00:00    36.380001
2015-01-02 02:00:00    36.380001
Freq: H, Name: Close, dtype: float64
```

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [21]: df_twtr_B.Close['2020-11-02':'2020-11-06']
```

```
Out[21]: Date
2020-11-02    39.470001
2020-11-03    41.730000
2020-11-04    42.759998
2020-11-05    43.709999
2020-11-06    43.119999
Freq: B, Name: Close, dtype: float64
```


Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [21]: df_twtr_B.Close['2020-11-02':'2020-11-06']
```

```
Out[21]: Date
2020-11-02    39.470001
2020-11-03    41.730000
2020-11-04    42.759998
2020-11-05    43.709999
2020-11-06    43.119999
Freq: B, Name: Close, dtype: float64
```

```
In [22]: rolling = df_twtr_B.Close.rolling(5, center=True)
rolling
```

```
Out[22]: Rolling [window=5,center=True,axis=0,method=single]
```

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [21]: df_twtr_B.Close['2020-11-02':'2020-11-06']
```

```
Out[21]: Date
2020-11-02    39.470001
2020-11-03    41.730000
2020-11-04    42.759998
2020-11-05    43.709999
2020-11-06    43.119999
Freq: B, Name: Close, dtype: float64
```

```
In [22]: rolling = df_twtr_B.Close.rolling(5, center=True)
rolling
```

```
Out[22]: Rolling [window=5,center=True,axis=0,method=single]
```

```
In [23]: rolling.mean()['2020-11-02':'2020-11-06']
```

```
Out[23]: Date
2020-11-02    43.550000
2020-11-03    41.806000
2020-11-04    42.157999
2020-11-05    42.901999
2020-11-06    43.037999
Freq: B, Name: Close, dtype: float64
```

Moving Windows

Moving Windows

```
In [24]: sns.set_style("whitegrid")
fig,ax = plt.subplots(1,1,figsize=(24,8));
df_twtr_B.loc['2020'].Close.plot(style='-',alpha=0.3,label='business day');
rolling.mean().loc['2020'].plot(style='--',label='5 day rolling window mean');
(rolling.mean().loc['2020'] + 2*rolling.std().loc['2020']).plot(style=':',c='g',label='_nolegend_');
(rolling.mean().loc['2020'] - 2*rolling.std().loc['2020']).plot(style=':',c='g',label='_nolegend_');
ax.legend();
```



Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [25]: df_bike = pd.read_csv('../data/FremontBridge_2015-2017.csv', parse_dates=['Date'], index_col='Date')
df_bike.columns = ['Total', 'East', 'West']
df_bike.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 52608 entries, 2015-01-01 00:00:00 to 2017-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total   52598 non-null     float64
1   East    52598 non-null     float64
2   West    52598 non-null     float64
dtypes: float64(3)
memory usage: 1.6 MB
```

Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [25]: df_bike = pd.read_csv('../data/FremontBridge_2015-2017.csv', parse_dates=['Date'], index_col='Date')
df_bike.columns = ['Total', 'East', 'West']
df_bike.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 52608 entries, 2015-01-01 00:00:00 to 2017-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total    52598 non-null     float64
1   East     52598 non-null     float64
2   West     52598 non-null     float64
dtypes: float64(3)
memory usage: 1.6 MB
```

```
In [26]: df_bike.head(3)
```

Out[26]:

	Total	East	West
Date			
2015-01-01 00:00:00	13.0	4.0	9.0
2015-01-01 01:00:00	27.0	4.0	23.0
2015-01-01 02:00:00	19.0	5.0	14.0

Example: Fill Missing Values

Example: Fill Missing Values

```
In [27]: f'proportion missing: {sum(df_bike.Total.isna()) / len(df_bike):0.5f}'
```

```
Out[27]: 'proportion missing: 0.00019'
```

Example: Fill Missing Values

```
In [27]: f'proportion missing: {sum(df_bike.Total.isna()) / len(df_bike):0.5f}'
```

```
Out[27]: 'proportion missing: 0.00019'
```

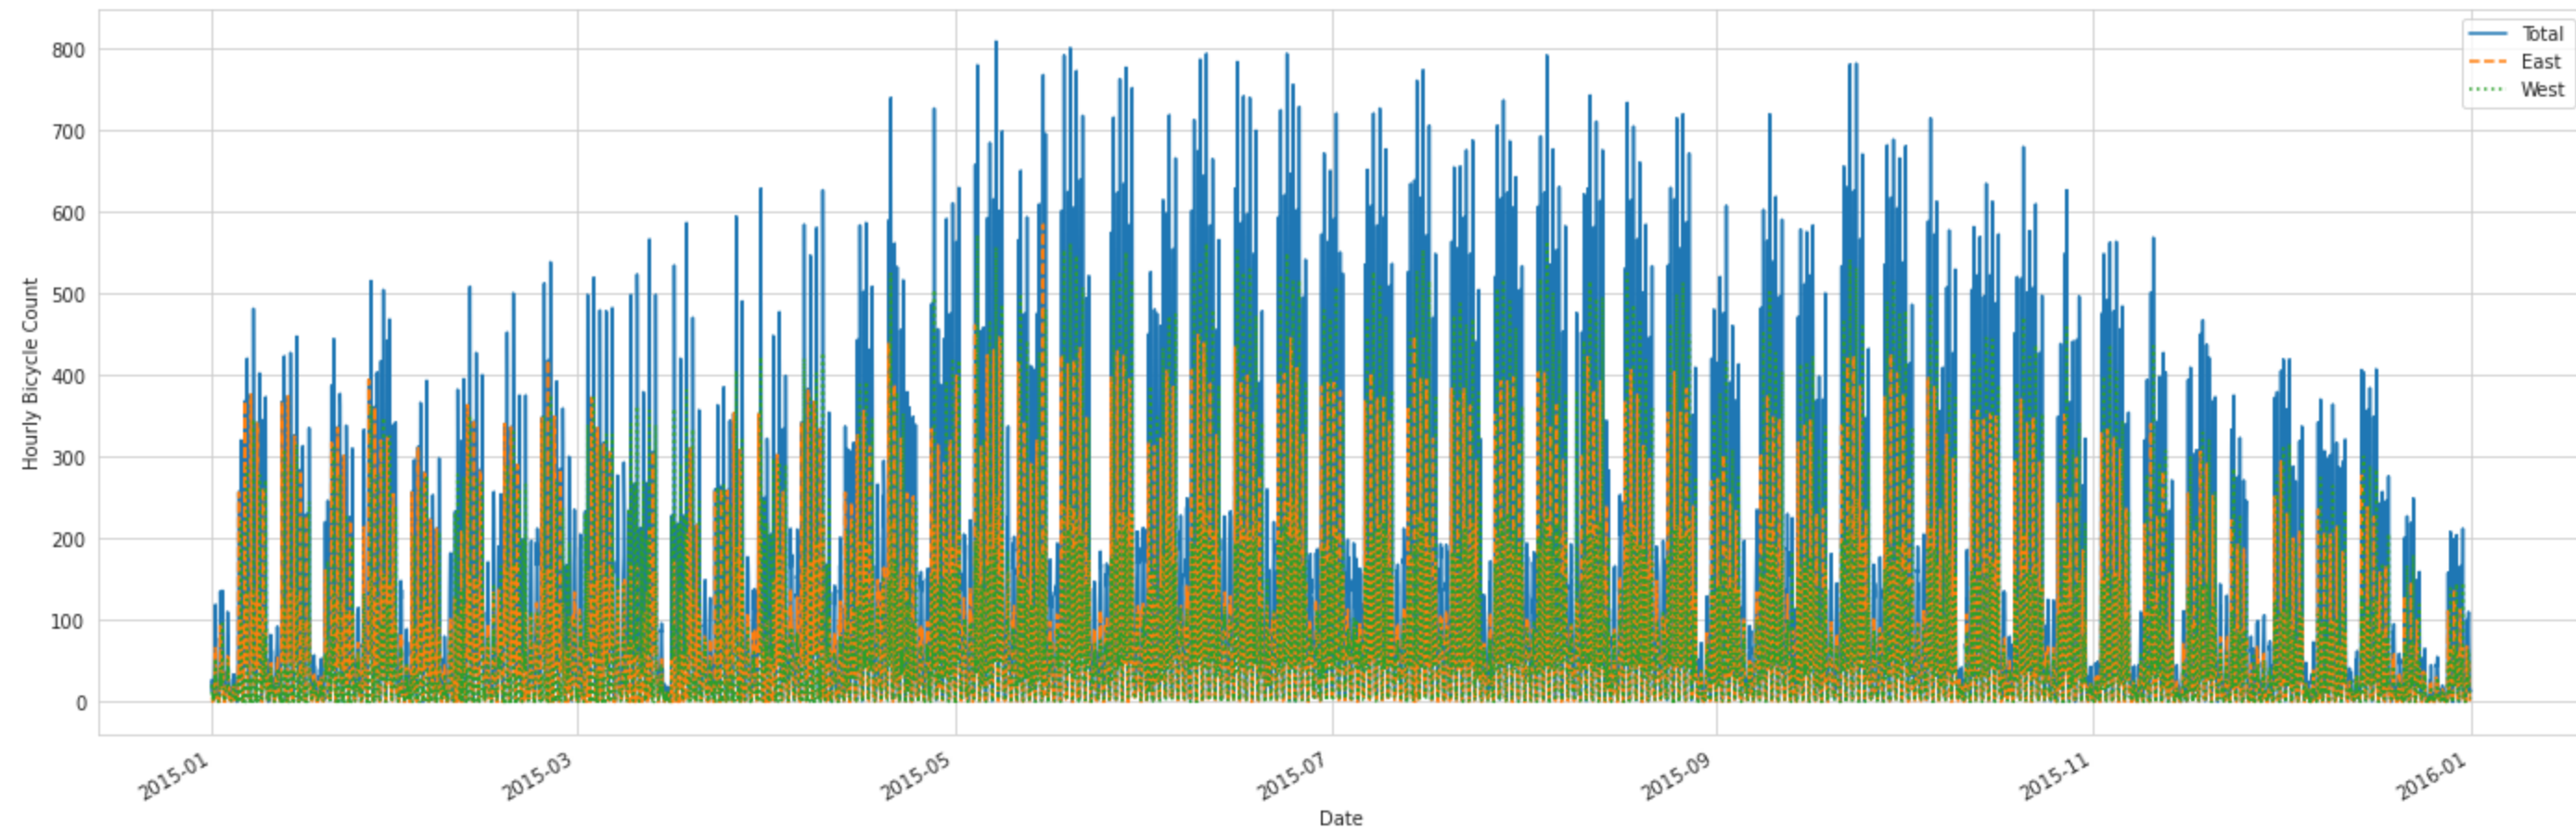
```
In [28]: df_bike = df_bike.fillna(method='ffill')
display(df_bike.head(3))
```

	Total	East	West
Date			
2015-01-01 00:00:00	13.0	4.0	9.0
2015-01-01 01:00:00	27.0	4.0	23.0
2015-01-01 02:00:00	19.0	5.0	14.0

Plot data from 2015

Plot data from 2015


```
In [29]: fig,ax = plt.subplots(1,1,figsize=(24,8))
df_bike.loc['2015'].plot(style=['-', '--', ':'],ax=ax)
plt.ylabel('Hourly Bicycle Count');
```



Downsample to weekly sum to smooth things out

Downsample to weekly sum to smooth things out


```
In [30]: weekly = df_bike.resample('W').sum()  
weekly.plot(style=[':', '--', '-'], figsize=(24,8))  
plt.ylabel('Weekly bicycle count');
```



Resample at daily for a more granular view and apply a rolling window of 30 days

Resample at daily for a more granular view and apply a rolling window of 30 days

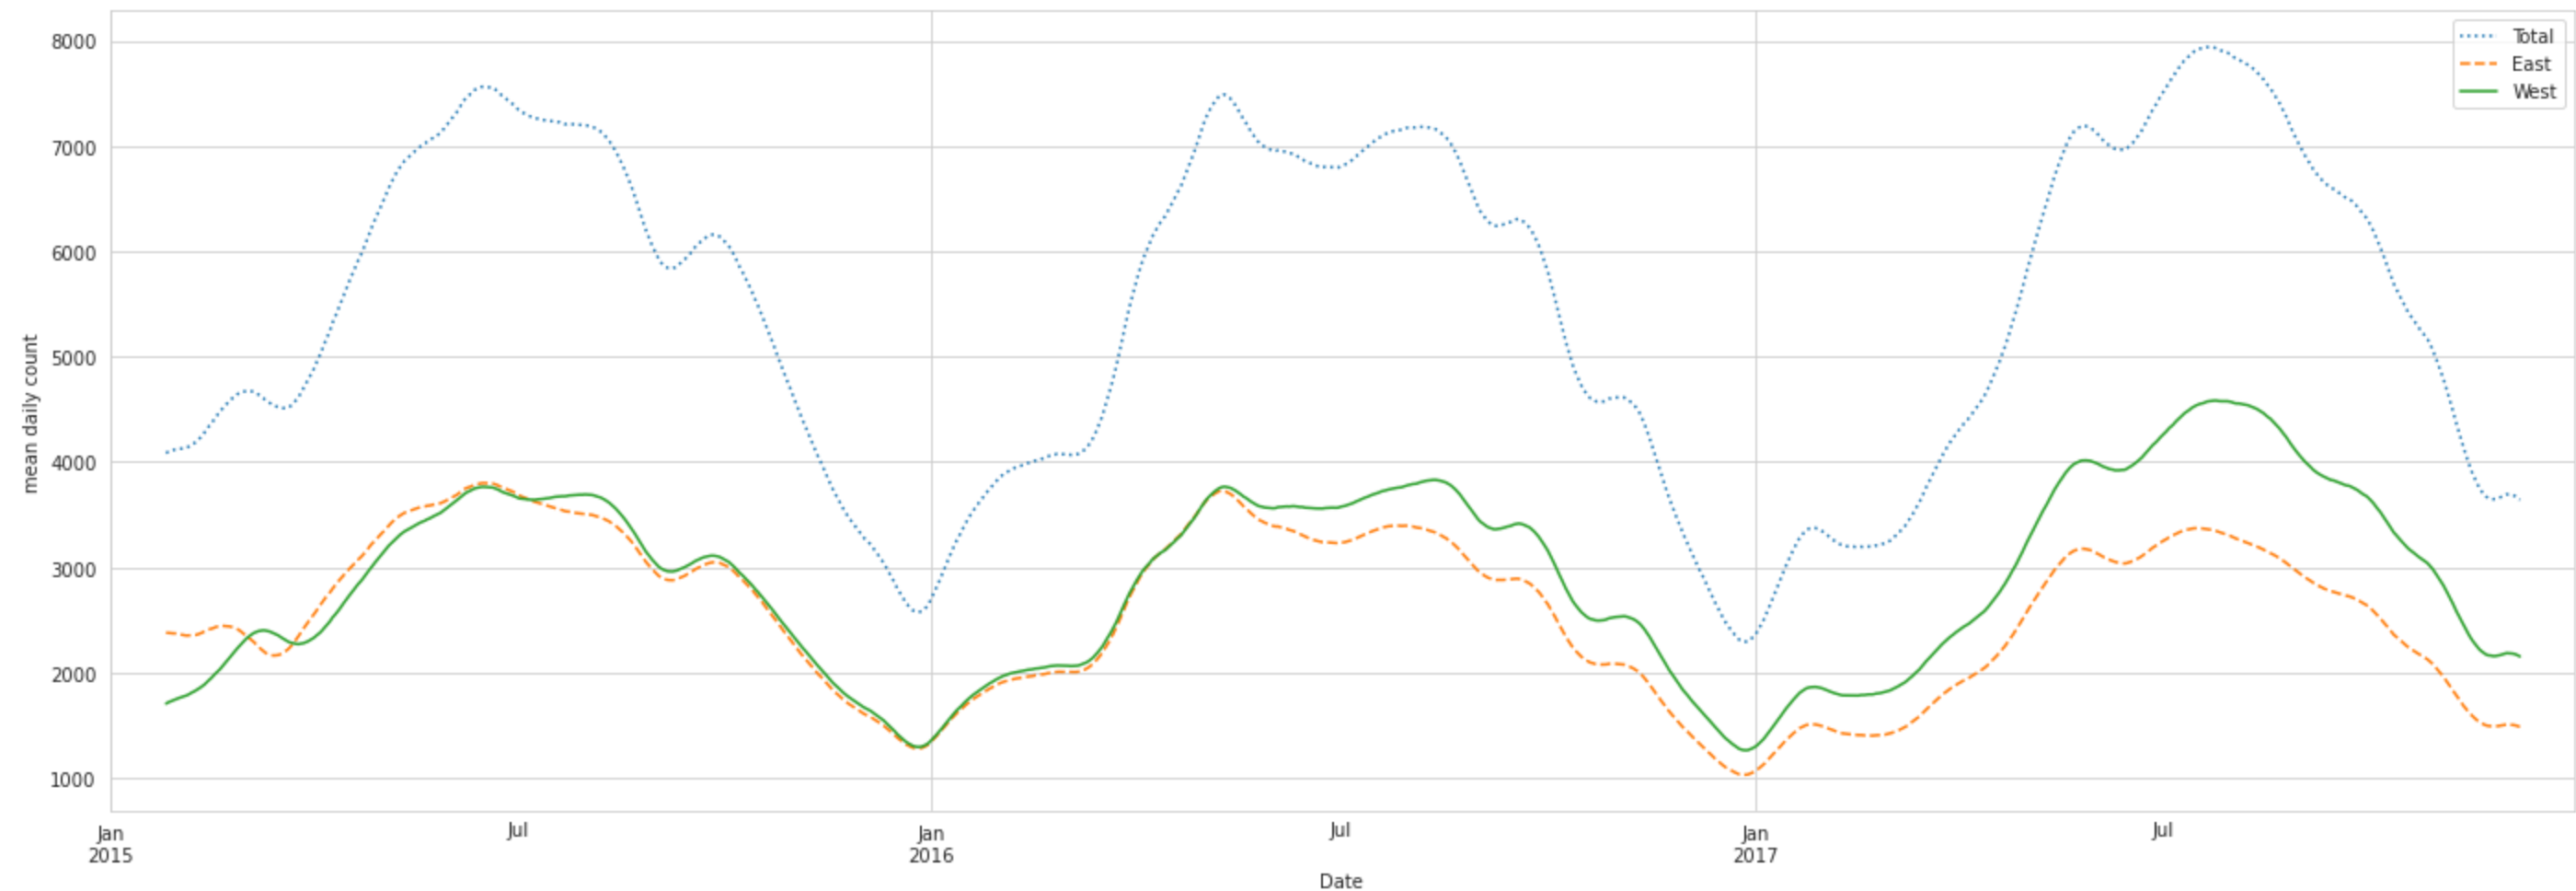
```
In [31]: daily = df_bike.resample('D').sum()  
daily.rolling(30,center=True).mean().plot(style=[':', '--', '-'], figsize=(24,8))  
plt.ylabel('mean daily count');
```



A wider window using a gaussian filter smooths more while accentuating daily differences

A wider window using a gaussian filter smooths more while accentuating daily differences

```
In [32]: daily.rolling(50,center=True,win_type='gaussian').mean(std=10).plot(style=[':', '--', '-'],figsize=(24,8));  
plt.ylabel('mean daily count');
```



From Datetime to Time

From Datetime to Time

```
In [33]: #If we want to only look at time of day  
df_bike.index.time
```

```
Out[33]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,  
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],  
              dtype=object)
```

From Datetime to Time

```
In [33]: #If we want to only look at time of day  
df_bike.index.time
```

```
Out[33]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,  
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],  
              dtype=object)
```

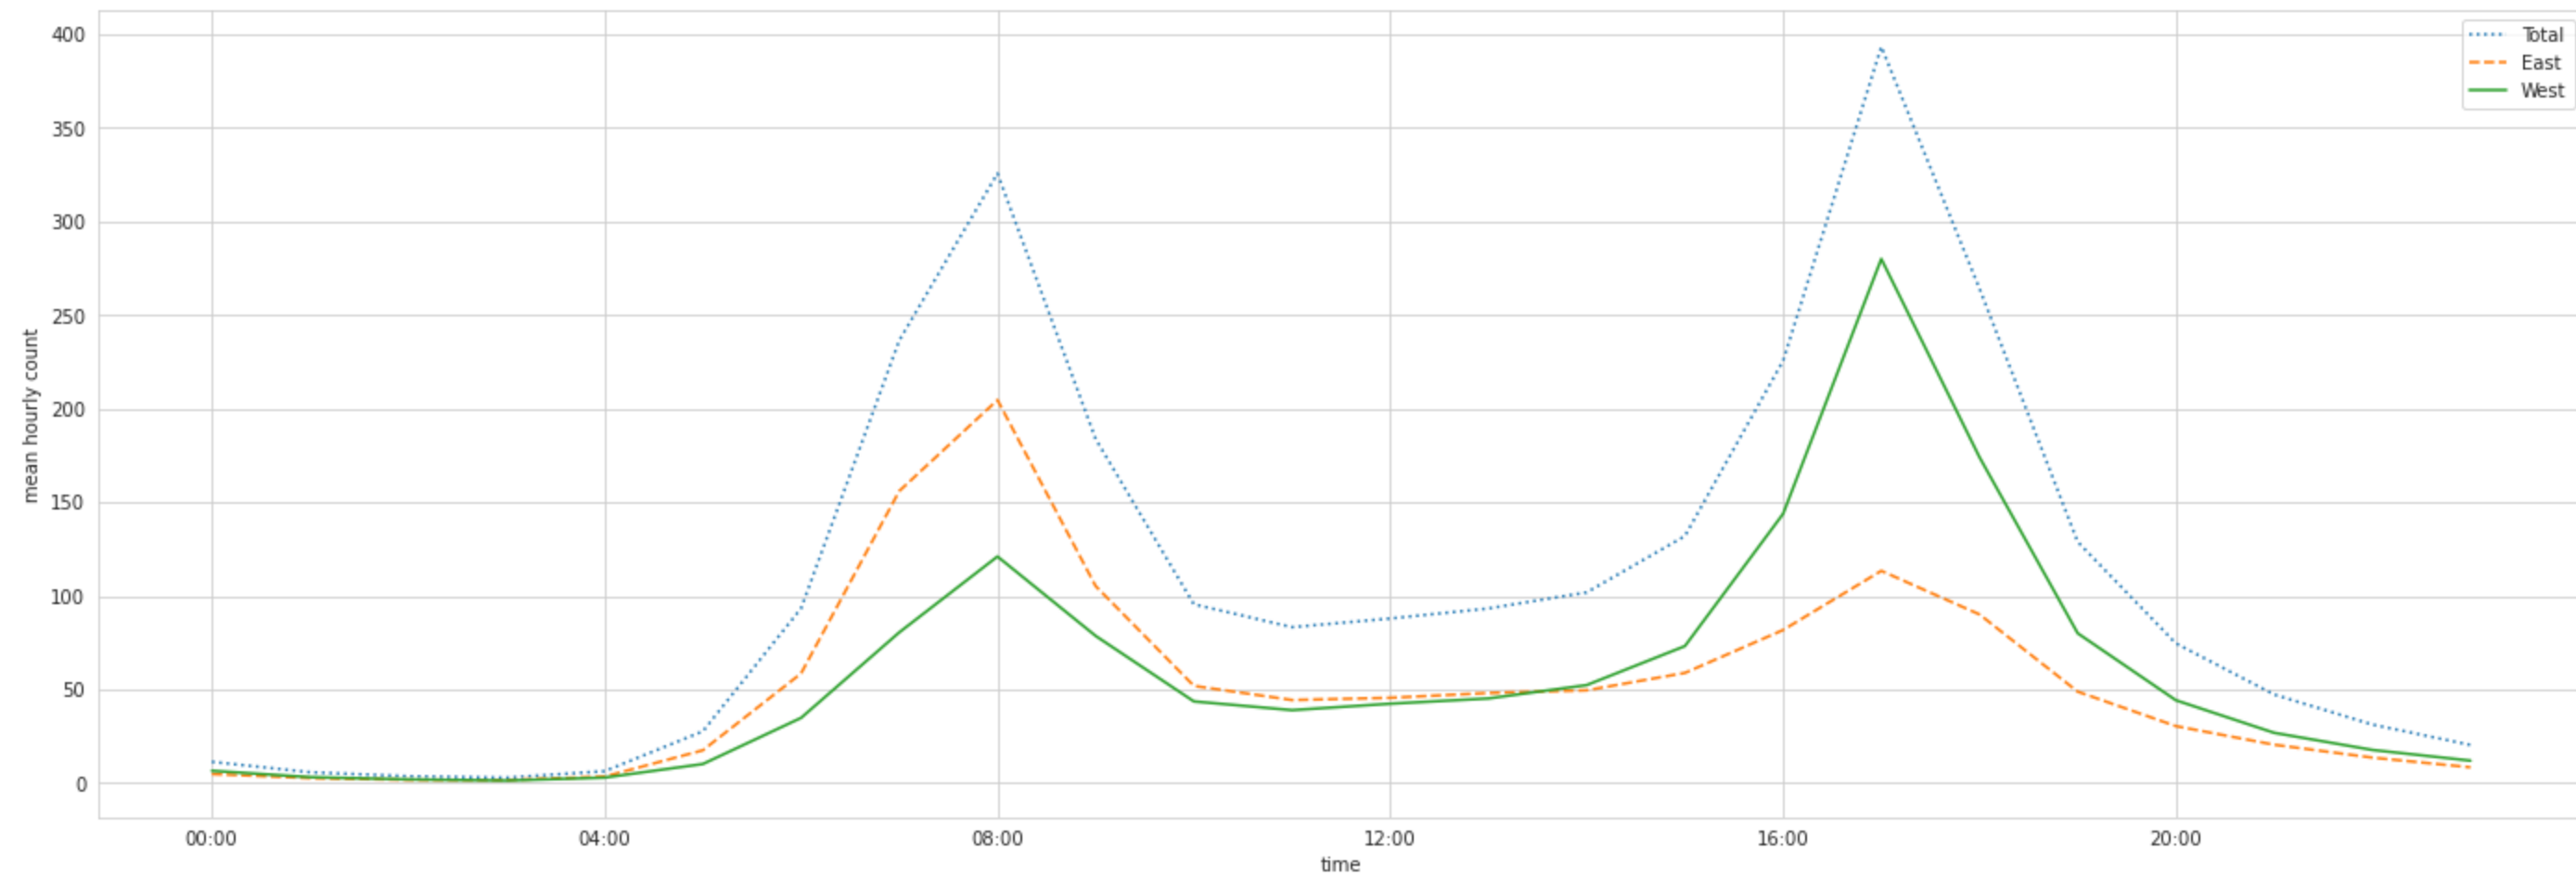
```
In [34]: #Get mean data by time (hourly)  
by_time = df_bike.groupby(df_bike.index.time).mean()  
display(by_time)
```

	Total	East	West
00:00:00	11.319343	4.806569	6.512774
01:00:00	5.743613	2.613139	3.130474
02:00:00	3.615876	1.687956	1.927920
03:00:00	2.850365	1.535584	1.314781
04:00:00	6.301095	3.418796	2.882299
05:00:00	27.629562	17.487226	10.142336
06:00:00	93.326642	58.534672	34.791971
07:00:00	236.456204	155.980839	80.475365
08:00:00	325.637774	204.628650	121.009124
09:00:00	183.750000	105.175182	78.574818
10:00:00	95.483577	51.891423	43.592153
11:00:00	83.261861	44.340328	38.921533
12:00:00	87.883212	45.524635	42.358577
13:00:00	93.259124	48.087591	45.171533
14:00:00	101.838504	49.528285	52.310219
15:00:00	131.947993	58.810219	73.137774
16:00:00	225.548358	81.741788	143.806569
17:00:00	393.354015	113.352190	280.001825
18:00:00	264.122263	90.069343	174.052920
19:00:00	128.802007	48.809307	79.992701

Plot by hour of the day

Plot by hour of the day

```
In [35]: hourly_ticks = 60 * 60 * 4 * np.arange(6) # sec * min * every4hours
by_time.plot(xticks=hourly_ticks, style=[':', '--', '-'], figsize=(24,8));
plt.ylabel('mean hourly count');
```

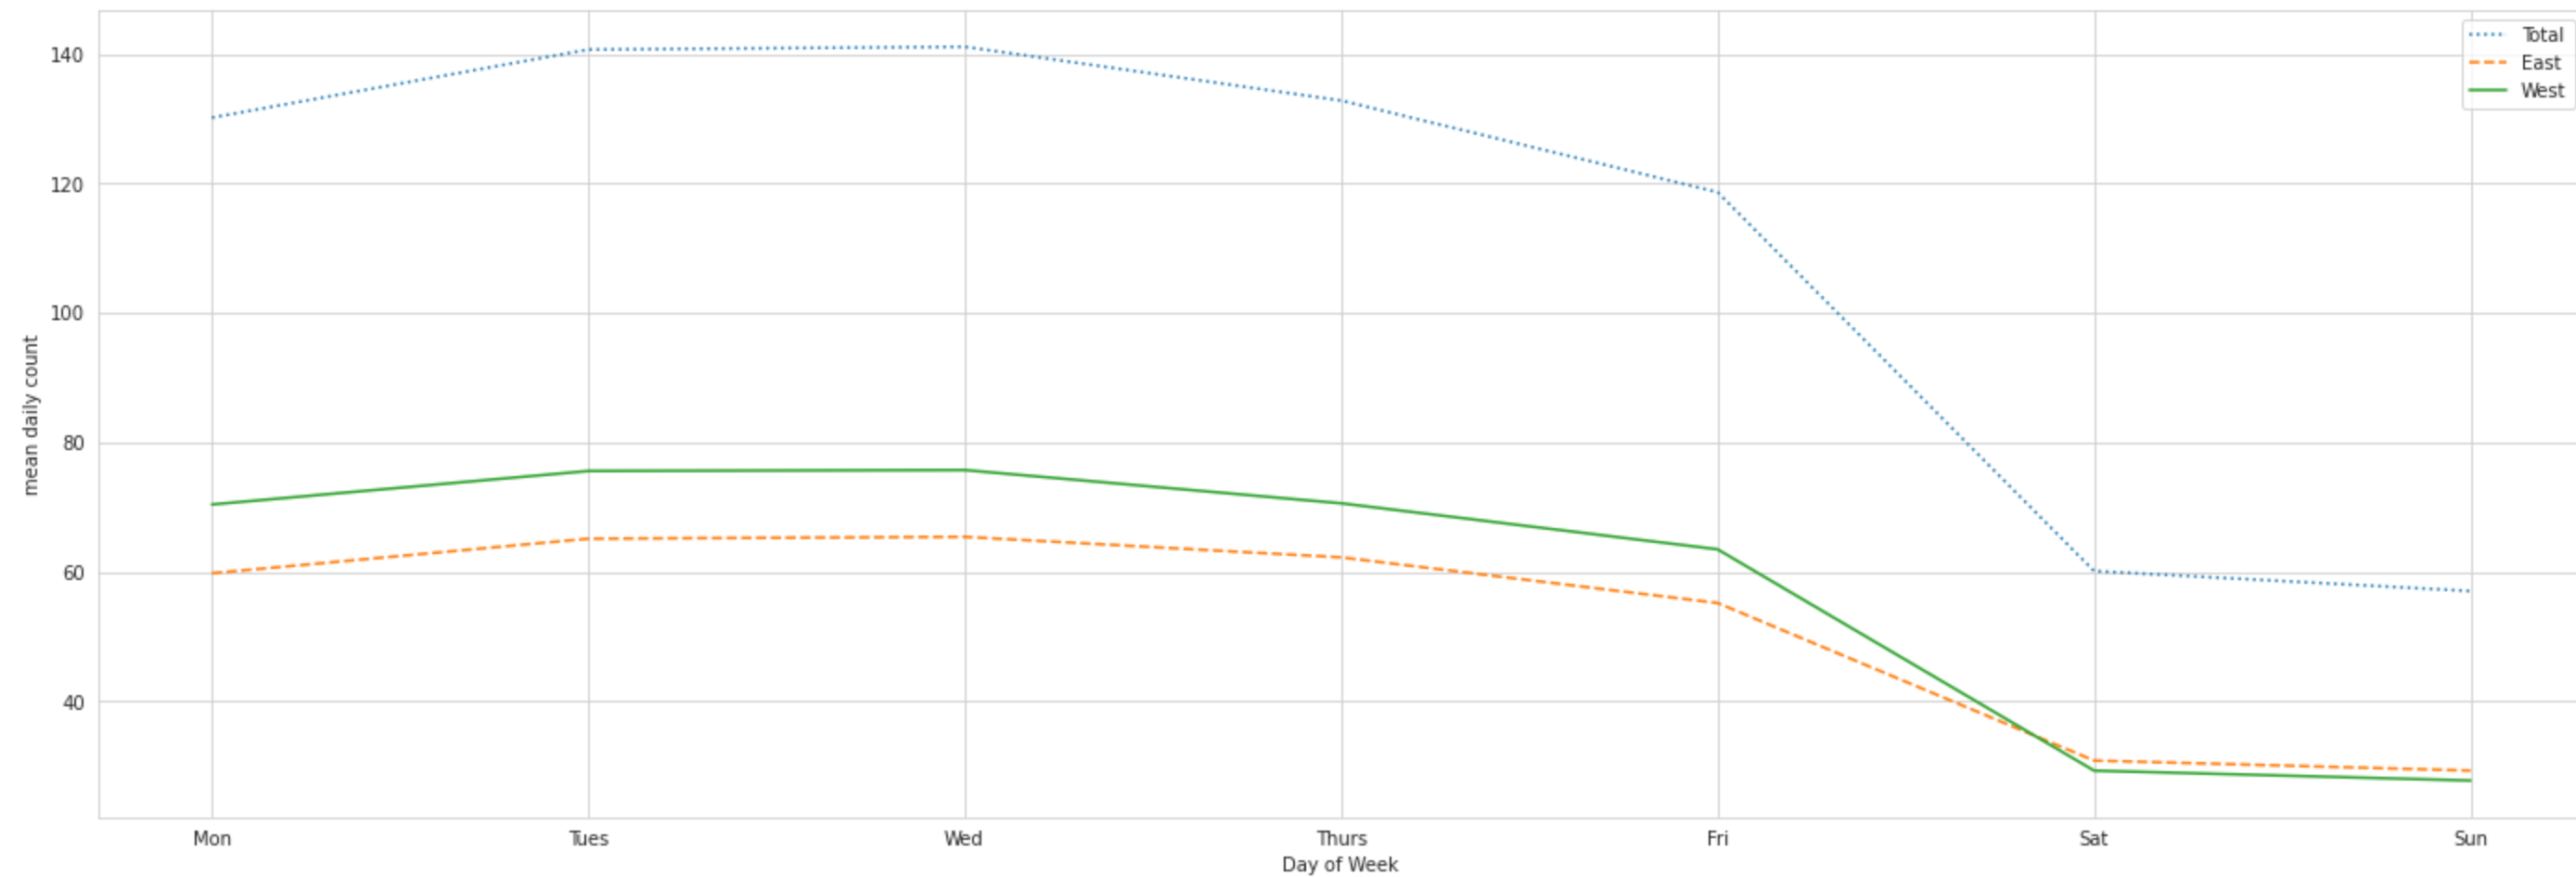


Can also look at average by day of week

Can also look at average by day of week

```
In [36]: # note that for dayofweek: 0 == Mon, 1 == Tues, ..., 6 == 'Sun'
by_weekday = df_bike.groupby(df_bike.index.dayofweek).mean()
by_weekday = by_weekday.set_index(pd.Index(['Mon', 'Tues', 'Wed', 'Thurs', 'Fri', 'Sat', 'Sun']))

fig, ax = plt.subplots(1, 1, figsize=(24, 8))
by_weekday.plot(style=[':', '--', '-'], ax=ax);
ax.set_xlabel('Day of Week'); ax.set_ylabel('mean daily count');
```

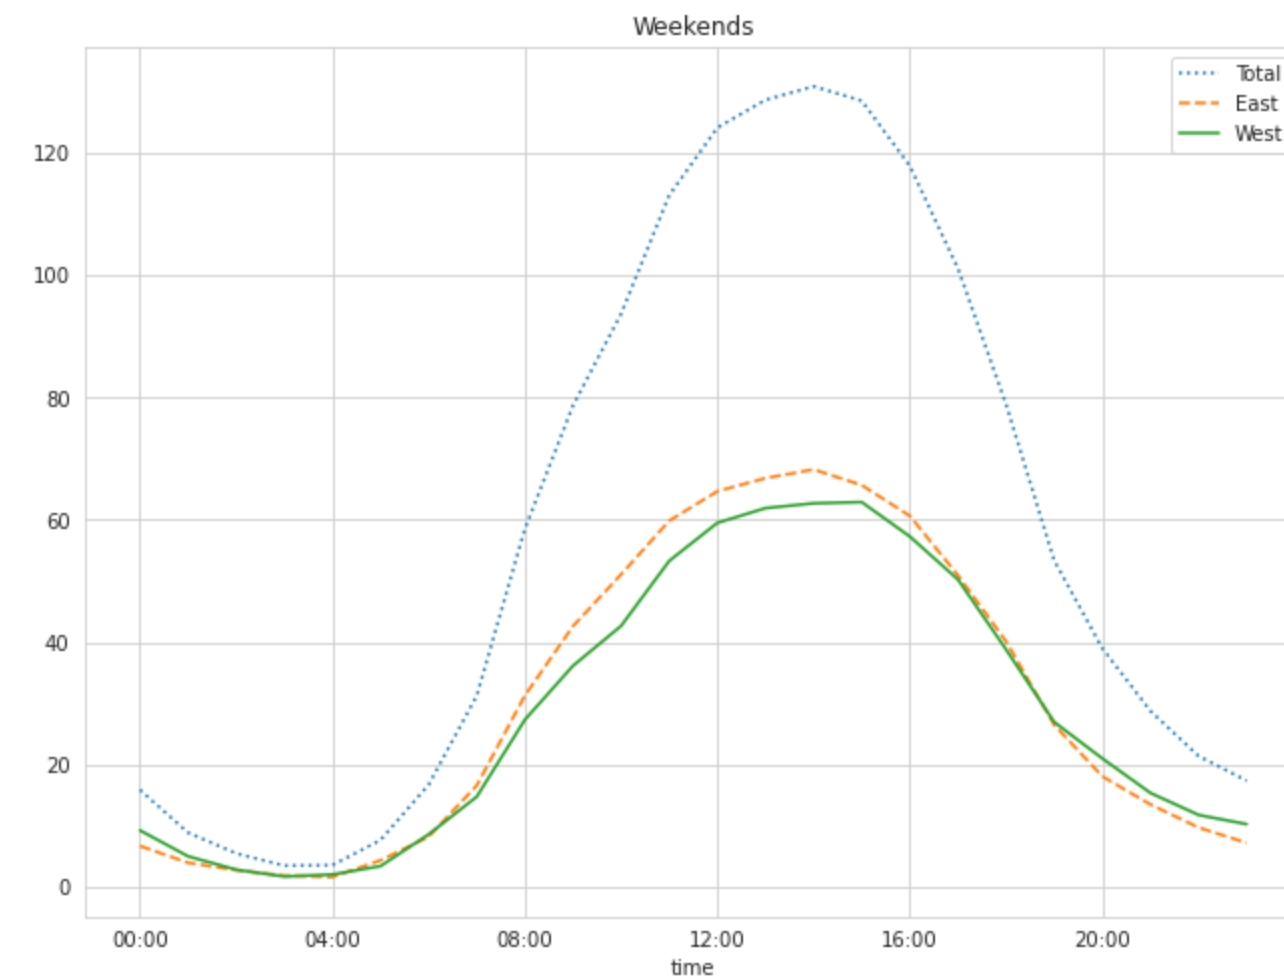
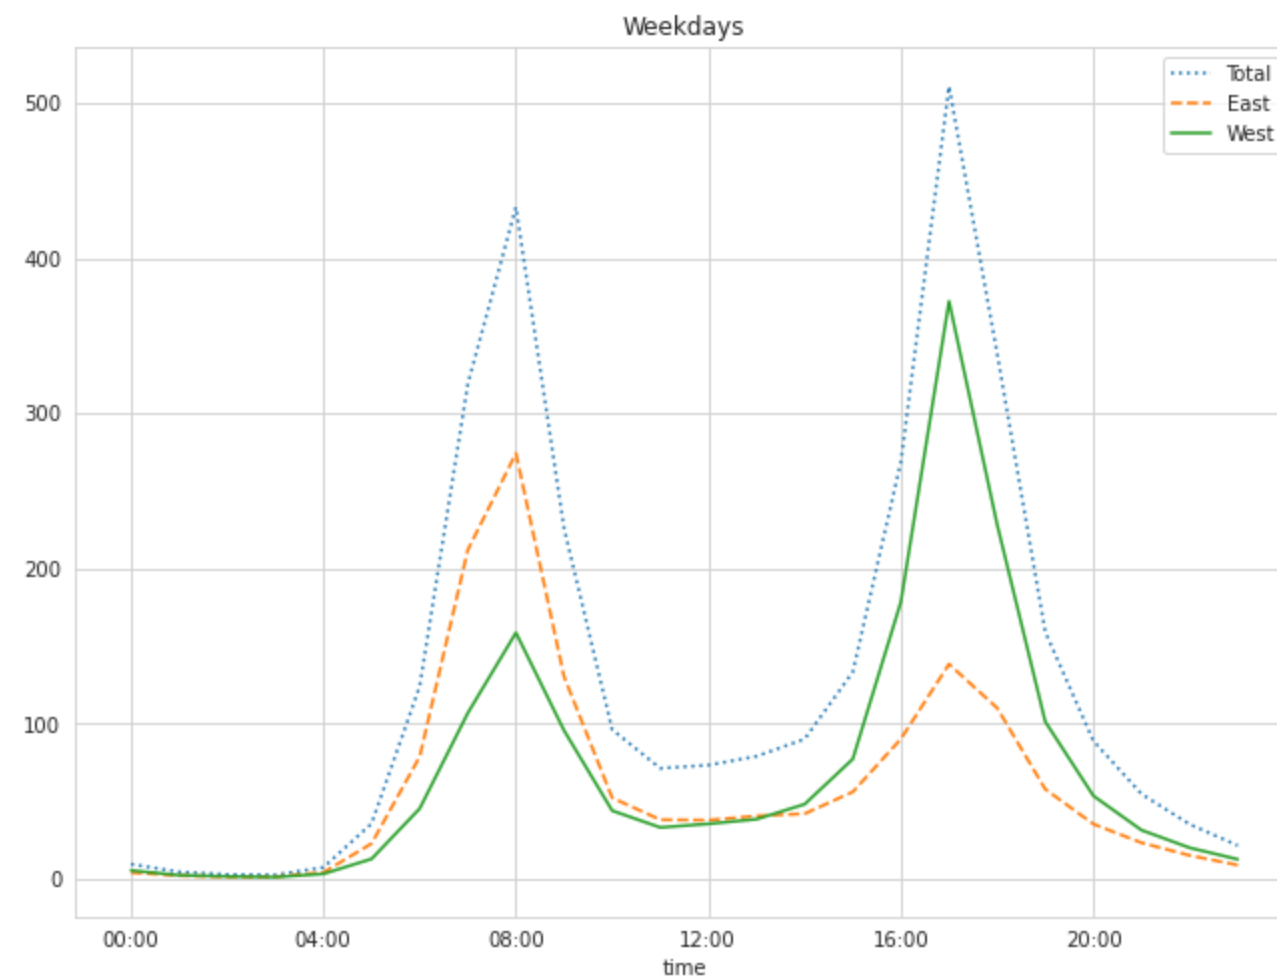


Separate out weekdays and weekends

Separate out weekdays and weekends

```
In [37]: # create a weekend mask
weekend = np.where(df_bike.index.weekday < 5, 'Weekday', 'Weekend')

# get hourly mean values split by weekday, weekend
by_time = df_bike.groupby([weekend, df_bike.index.time]).mean()
fig, ax = plt.subplots(1, 2, figsize=(24, 8))
by_time.loc['Weekday'].plot(ax=ax[0], title='Weekdays', xticks=hourly_ticks, style=[':', '--', '-'])
by_time.loc['Weekend'].plot(ax=ax[1], title='Weekends', xticks=hourly_ticks, style=[':', '--', '-']);
```



Can we predict daily Total bike traffic?

Can we predict daily Total bike traffic?

```
In [38]: bike_counts = pd.read_csv('../data/FremontBridge.csv', index_col='Date', parse_dates=True)
bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)

bike_counts = bike_counts.loc[:, ['Fremont Bridge Total']] # Keep Total as target
bike_counts.columns = ['Total']
daily_data = bike_counts.resample('d').sum() # downsample to daily totals
print(daily_data.head(3))
```

Date	Total
2012-10-03	7042.0
2012-10-04	6950.0
2012-10-05	6296.0

Can we predict daily Total bike traffic?

```
In [38]: bike_counts = pd.read_csv('../data/FremontBridge.csv', index_col='Date', parse_dates=True)
bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)

bike_counts = bike_counts.loc[:, ['Fremont Bridge Total']] # Keep Total as target
bike_counts.columns = ['Total']
daily_data = bike_counts.resample('d').sum() # downsample to daily totals
print(daily_data.head(3))
```

Date	Total
2012-10-03	7042.0
2012-10-04	6950.0
2012-10-05	6296.0

On to Feature Engineering...

Add 'day of week' one-hot features

Add 'day of week' one-hot features

```
In [39]: days = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']
         for i in range(7):
             daily_data[days[i]] = (daily_data.index.dayofweek == i).astype(float)
         print(daily_data.head(3))
```

	Total	Mon	Tue	Wed	Thu	Fri	Sat	Sun
Date								
2012-10-03	7042.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
2012-10-04	6950.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
2012-10-05	6296.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0

Add 'is it a holiday' dummy feature

Add 'is it a holiday' dummy feature

```
In [40]: from pandas.tseries.holiday import USFederalHolidayCalendar
cal = USFederalHolidayCalendar()
holidays = cal.holidays('2012', '2016')

daily_data = daily_data.join(pd.Series(1, index=holidays, name='holiday'))
daily_data['holiday'].fillna(0, inplace=True)
print(daily_data.head(3))
```

	Total	Mon	Tue	Wed	Thu	Fri	Sat	Sun	holiday
Date									
2012-10-03	7042.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
2012-10-04	6950.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
2012-10-05	6296.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

Add number of hours of daylight

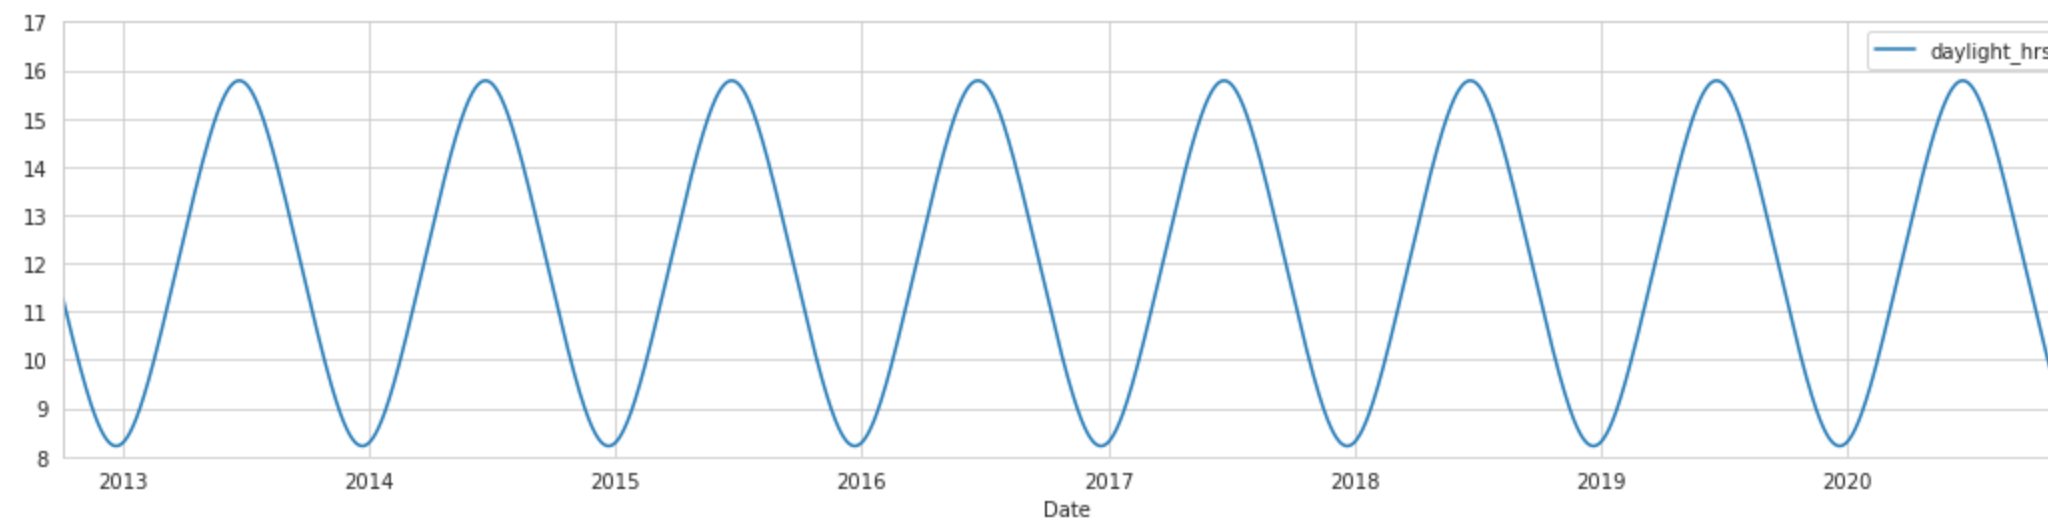
Add number of hours of daylight

```
In [41]: from datetime import datetime

def hours_of_daylight(date, axis=23.44, latitude=47.61):
    """Compute the hours of daylight for the given date"""
    days = (date - datetime(2000, 12, 21)).days
    m = (1. - np.tan(np.radians(latitude))
         * np.tan(np.radians(axis) * np.cos(days * 2 * np.pi / 365.25)))
    return 24. * np.degrees(np.arccos(1 - np.clip(m, 0, 2))) / 180.

daily_data['daylight_hrs'] = list(map(hours_of_daylight, daily_data.index));

daily_data[['daylight_hrs']].plot(figsize=(18,4));
plt.ylim(8, 17);
```



Add weather information (can we predict this for future dates?)

Add weather information (can we predict this for future dates?)

```
In [42]: # temperatures are in 1/10 deg C; convert to C
bike_weather['TMIN'] /= 10
bike_weather['TMAX'] /= 10
bike_weather['Temp (C)'] = 0.5 * (bike_weather['TMIN'] + bike_weather['TMAX'])

# precip is in 1/10 mm; convert to inches
bike_weather['PRCP'] /= 254
bike_weather['dry day'] = (bike_weather['PRCP'] == 0).astype(int)

daily_data = daily_data.join(bike_weather[['PRCP', 'Temp (C)', 'dry day']])
daily_data.head(3)
```

Out[42]:

	Total	Mon	Tue	Wed	Thu	Fri	Sat	Sun	holiday	daylight_hrs	PRCP	Temp (C)	dry day
Date													
2012-10-03	7042.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	11.277359	0.0	13.35	1.0
2012-10-04	6950.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	11.219142	0.0	13.60	1.0
2012-10-05	6296.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	11.161038	0.0	15.30	1.0

Add time of year

Add time of year

```
In [43]: daily_data['annual'] = (daily_data.index - daily_data.index[0]).days / 365.  
daily_data.head(3)
```

Out[43]:

	Total	Mon	Tue	Wed	Thu	Fri	Sat	Sun	holiday	daylight_hrs	PRCP	Temp (C)	dry day	annual
Date														
2012-10-03	7042.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	11.277359	0.0	13.35	1.0	0.000000
2012-10-04	6950.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	11.219142	0.0	13.60	1.0	0.002740
2012-10-05	6296.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	11.161038	0.0	15.30	1.0	0.005479

Generate a model

Generate a model

```
In [44]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# drop any rows with missing data
daily_data.dropna(axis=0, how='any', inplace=True)

X_bike = daily_data[daily_data.columns[daily_data.columns != 'Total']]
y_bike = daily_data.Total

X_bike_train, X_bike_test, y_bike_train, y_bike_test = train_test_split(X_bike, y_bike)

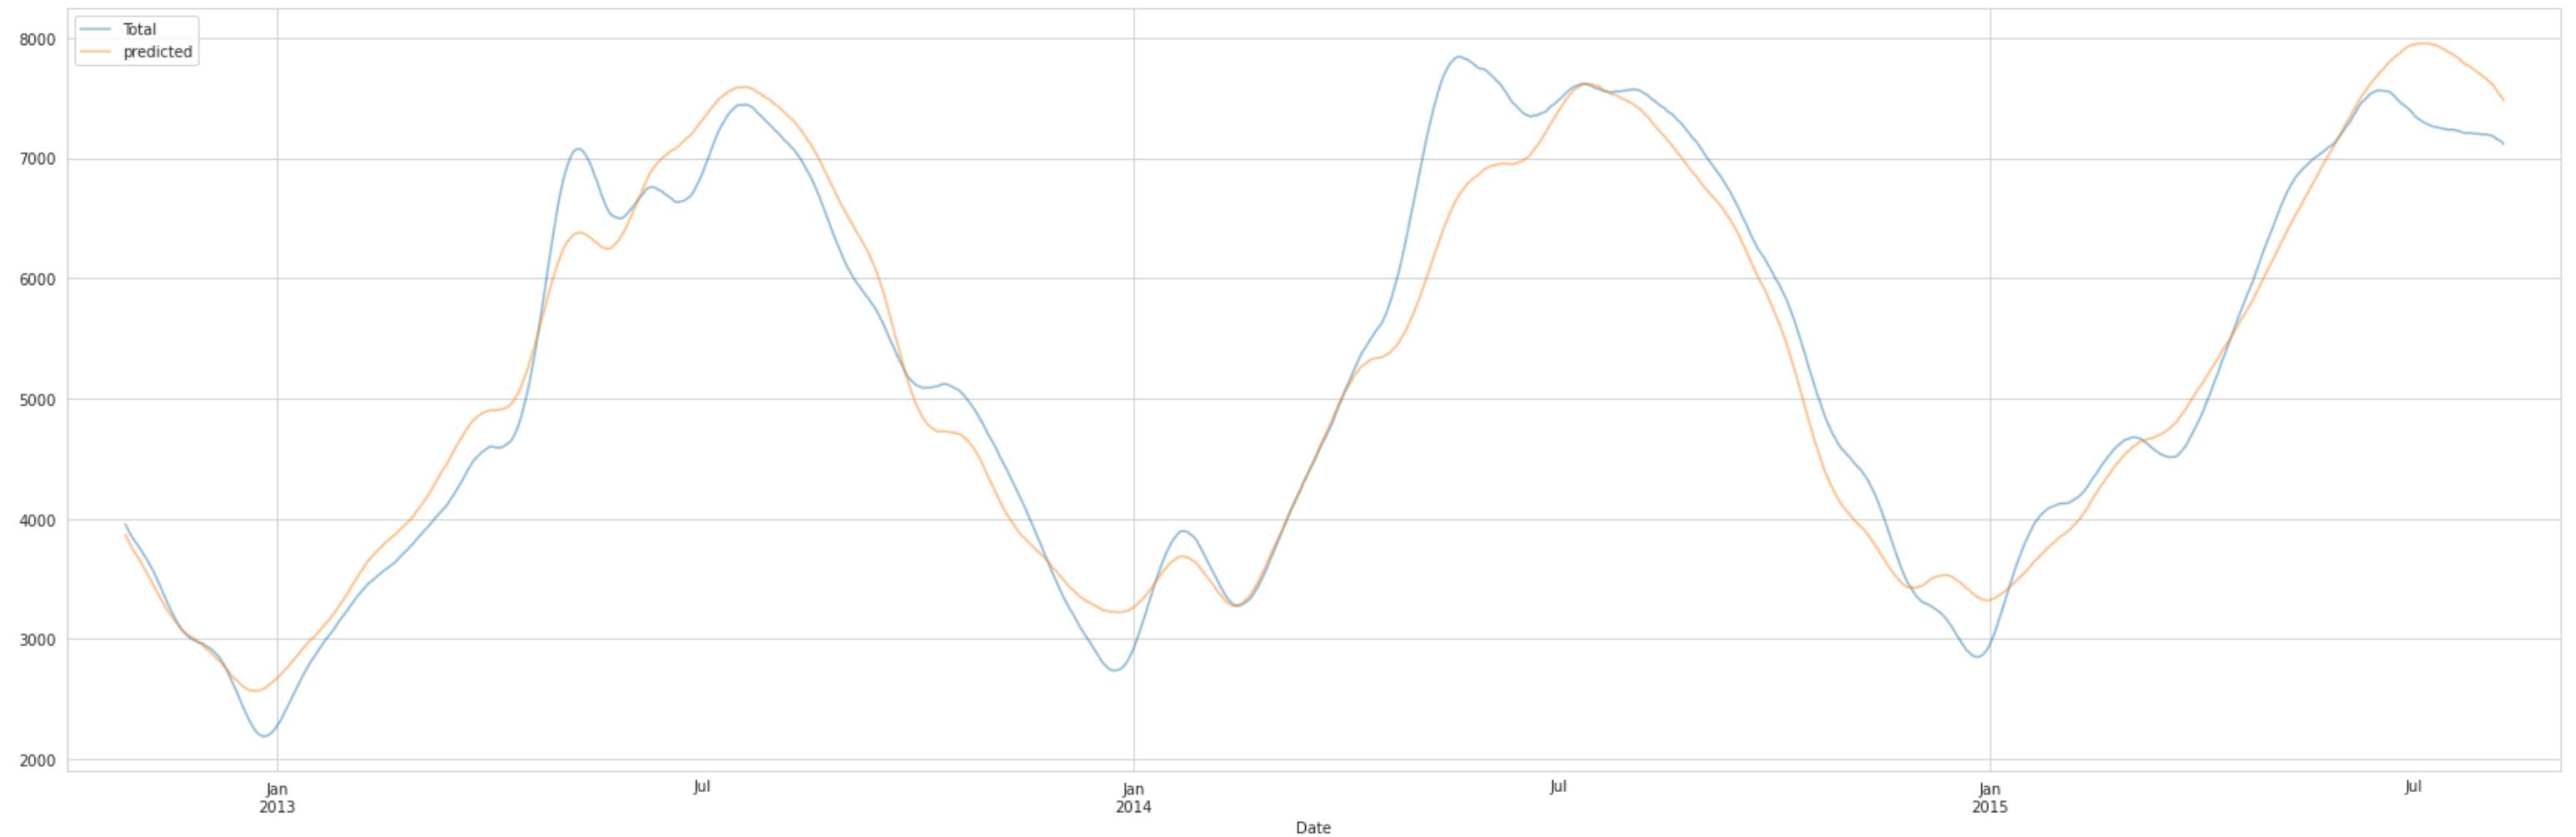
model = LinearRegression(fit_intercept=False)
model.fit(X_bike_train, y_bike_train)
print(f'training set R^2 : {model.score(X_bike_train, y_bike_train):0.2f}')
print(f'test set R^2      : {model.score(X_bike_test, y_bike_test):0.2f}')

training set R^2 : 0.87
test set R^2      : 0.86
```

Plot predictions vs observed

Plot predictions vs observed

```
In [45]: daily_data['predicted'] = model.predict(X_bike)
daily_data[['Total', 'predicted']].rolling(50, center=True, win_type='gaussian').mean(std=10).plot(alpha=0.5, figsize=(24, 8))
plt.tight_layout()
```



Time Series Operations Review

- Shifting
- Resampling
 - Downsampling
 - Upsampling
- Moving/Rolling Windows

Questions re Time Series Transformations?

Imbalanced Classes

- Imbalanced classes:
 - when there is significantly more of one class than another in a classification task
- common in real world datasets
- Ex: credit card fraud
 - very small number of fraud transactions relative to total transactions

Dealing With Imbalanced Classes

- Stratified Sampling
- Random Undersampling
- Random Oversampling
- Oversample Synthetic Minority Items
 - SMOTE
 - ADASYN
- Other methods

Stratified Sampling

Stratified Sampling

```
In [46]: from sklearn.model_selection import StratifiedKFold
```

```
X = np.ones(10)
```

```
y = [0, 0, 0, 0, 1, 1, 1, 1, 1, 1]
```

```
skf = StratifiedKFold(n_splits=3)
```

```
for train, test in skf.split(X, y):  
    print("%s %s" % (train, test))
```

```
[2 3 6 7 8 9] [0 1 4 5]
```

```
[0 1 3 4 5 8 9] [2 6 7]
```

```
[0 1 2 4 5 6 7] [3 8 9]
```

Random Sampling

- Randomly Oversample minority class
- Randomly Undersample majority class

Example Dataset

Example Dataset

```
In [47]: from sklearn.datasets import make_classification
from collections import Counter
X_imb, y_imb = make_classification(n_samples=5000, n_features=2, n_informative=2,
                                n_redundant=0, n_repeated=0, n_classes=3,
                                n_clusters_per_class=1,
                                weights=[0.01, 0.05, 0.94],
                                class_sep=0.8, random_state=0)

Counter(y_imb).items()
```

```
Out[47]: dict_items([(2, 4674), (1, 262), (0, 64)])
```

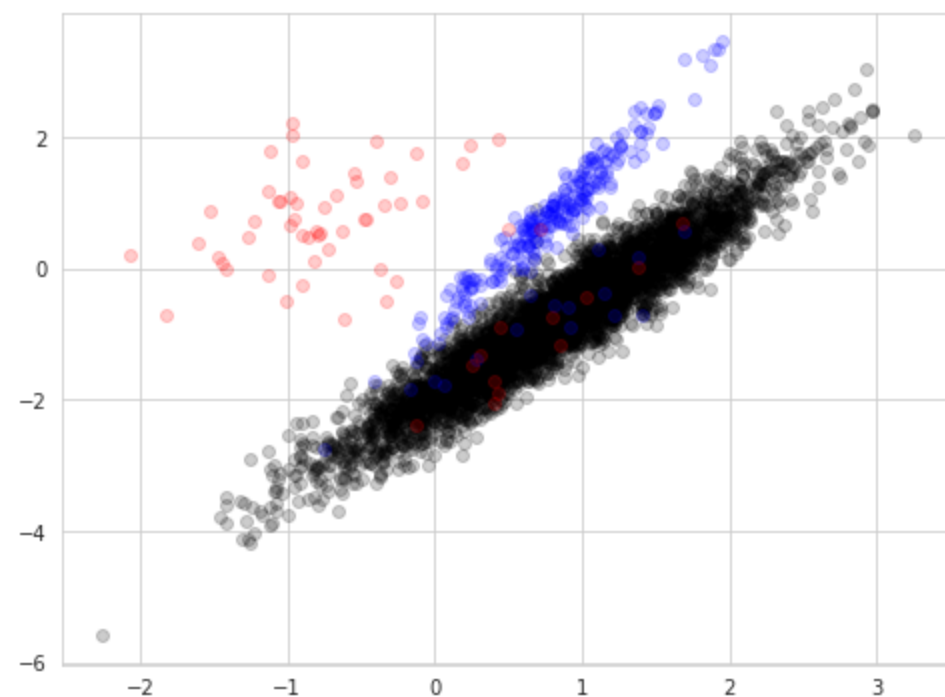
Example Dataset

```
In [47]: from sklearn.datasets import make_classification
from collections import Counter
X_imb, y_imb = make_classification(n_samples=5000, n_features=2, n_informative=2,
                                   n_redundant=0, n_repeated=0, n_classes=3,
                                   n_clusters_per_class=1,
                                   weights=[0.01, 0.05, 0.94],
                                   class_sep=0.8, random_state=0)

Counter(y_imb).items()
```

```
Out[47]: dict_items([(2, 4674), (1, 262), (0, 64)])
```

```
In [48]: fig, ax = plt.subplots(1, 1, figsize=(8, 6))
ax.scatter(X_imb[y_imb==2, 0], X_imb[y_imb==2, 1], c='k', alpha=.2);
ax.scatter(X_imb[y_imb==1, 0], X_imb[y_imb==1, 1], c='b', alpha=.2);
ax.scatter(X_imb[y_imb==0, 0], X_imb[y_imb==0, 1], c='r', alpha=.2);
```



Using imblearn

- `imblearn` is library to created to deal with imbalanced classes
- need to install from `conda-forge` as `imbalanced-learn`
- import from `imblearn`

Random Oversampling of minority class

Random Oversampling of minority class

```
In [49]: from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=0)  
X_ros, y_ros = ros.fit_resample(X_imb, y_imb)  
Counter(y_ros).items()
```

```
Out[49]: dict_items([(2, 4674), (1, 4674), (0, 4674)])
```

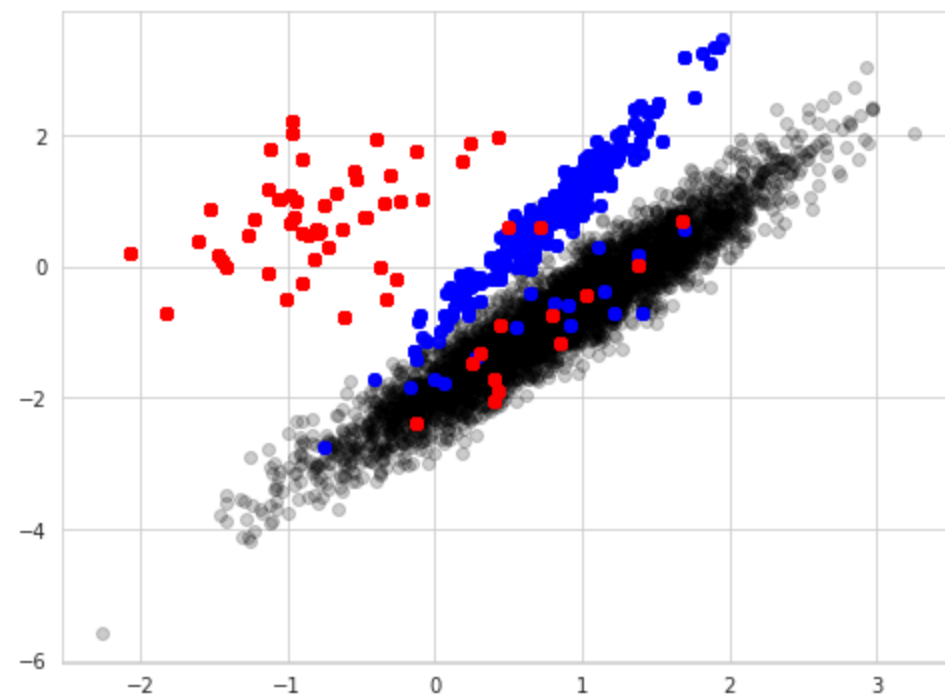
Random Oversampling of minority class

```
In [49]: from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=0)  
X_ros, y_ros = ros.fit_resample(X_imb, y_imb)  
Counter(y_ros).items()
```

```
Out[49]: dict_items([(2, 4674), (1, 4674), (0, 4674)])
```

```
In [50]: fig, ax = plt.subplots(1, 1, figsize=(8, 6))  
ax.scatter(X_ros[y_ros==2, 0], X_ros[y_ros==2, 1], c='k', alpha=.2);  
ax.scatter(X_ros[y_ros==1, 0], X_ros[y_ros==1, 1], c='b', alpha=.2);  
ax.scatter(X_ros[y_ros==0, 0], X_ros[y_ros==0, 1], c='r', alpha=.2);
```



Random Undersampling of majority class

Random Undersampling of majority class

```
In [51]: from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=0)  
X_rus, y_rus, = rus.fit_resample(X_imb, y_imb)  
Counter(y_rus).items()
```

```
Out[51]: dict_items([(0, 64), (1, 64), (2, 64)])
```

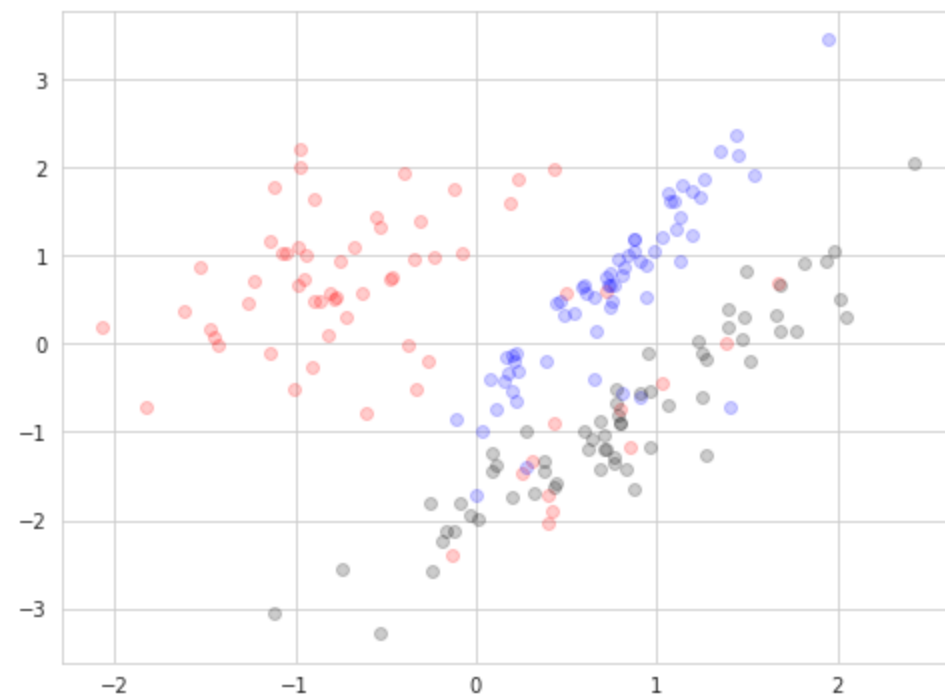
Random Undersampling of majority class

```
In [51]: from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=0)
X_rus, y_rus, = rus.fit_resample(X_imb, y_imb)
Counter(y_rus).items()
```

```
Out[51]: dict_items([(0, 64), (1, 64), (2, 64)])
```

```
In [52]: fig,ax=plt.subplots(1,1,figsize=(8,6))
ax.scatter(X_rus[y_rus==0,0],X_rus[y_rus==0,1],c='r', alpha=.2);
ax.scatter(X_rus[y_rus==1,0],X_rus[y_rus==1,1],c='b', alpha=.2);
ax.scatter(X_rus[y_rus==2,0],X_rus[y_rus==2,1],c='k', alpha=.2);
```



Oversample Sythetic Minority Items

- SMOTE: Synthetic Minority Oversampling
- ADASYN: Adaptive Synthetic Minority Oversampling

SMOTE: Synthetic Minority Oversampling

- Create new synthetic points between existing points

SMOTE: Synthetic Minority Oversampling

- Create new synthetic points between existing points

```
In [53]: from imblearn.over_sampling import SMOTE  
  
X_smote, y_smote = SMOTE().fit_resample(X_imb, y_imb)  
Counter(y_smote).items()
```

```
Out[53]: dict_items([(2, 4674), (1, 4674), (0, 4674)])
```


SMOTE: Synthetic Minority Oversampling

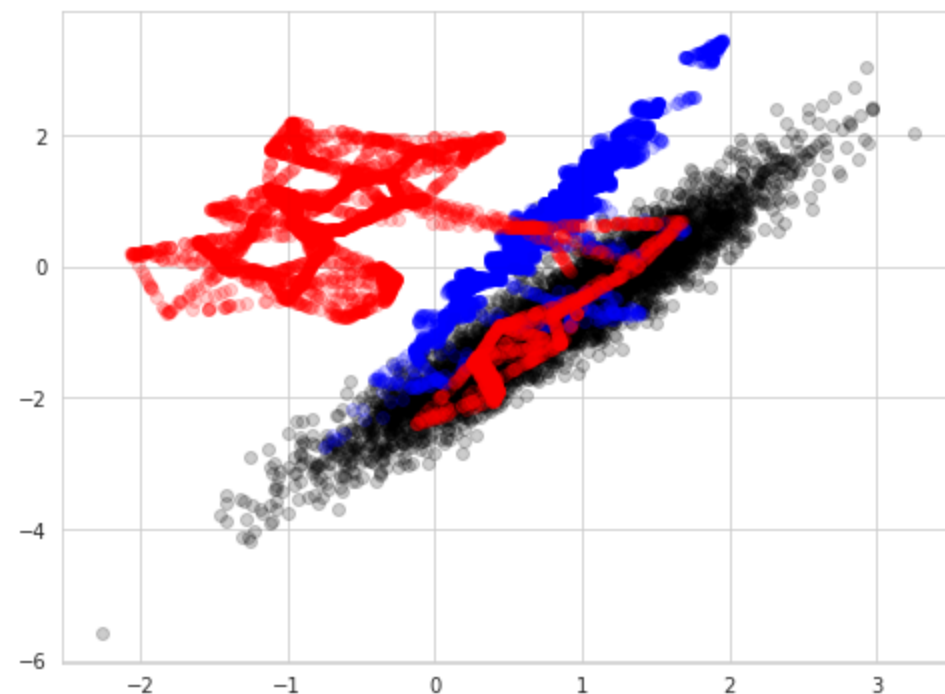
- Create new synthetic points between existing points

```
In [53]: from imblearn.over_sampling import SMOTE

X_smote, y_smote = SMOTE().fit_resample(X_imb, y_imb)
Counter(y_smote).items()
```

```
Out[53]: dict_items([(2, 4674), (1, 4674), (0, 4674)])
```

```
In [54]: fig,ax=plt.subplots(1,1,figsize=(8,6))
ax.scatter(X_smote[y_smote==2,0],X_smote[y_smote==2,1],c='k', alpha=.2);
ax.scatter(X_smote[y_smote==1,0],X_smote[y_smote==1,1],c='b', alpha=.2);
ax.scatter(X_smote[y_smote==0,0],X_smote[y_smote==0,1],c='r', alpha=.2);
```



ADASYN: Adaptive Synthetic Minority Oversampling

- Create new synthetic points between existing points *where classes overlap*

ADASYN: Adaptive Synthetic Minority Oversampling

- Create new synthetic points between existing points *where classes overlap*

```
In [55]: from imblearn.over_sampling import ADASYN  
  
X_adasyn, y_adasyn = ADASYN().fit_resample(X_imb, y_imb)  
Counter(y_adasyn).items()
```

```
Out[55]: dict_items([(2, 4674), (1, 4662), (0, 4673)])
```

ADASYN: Adaptive Synthetic Minority Oversampling

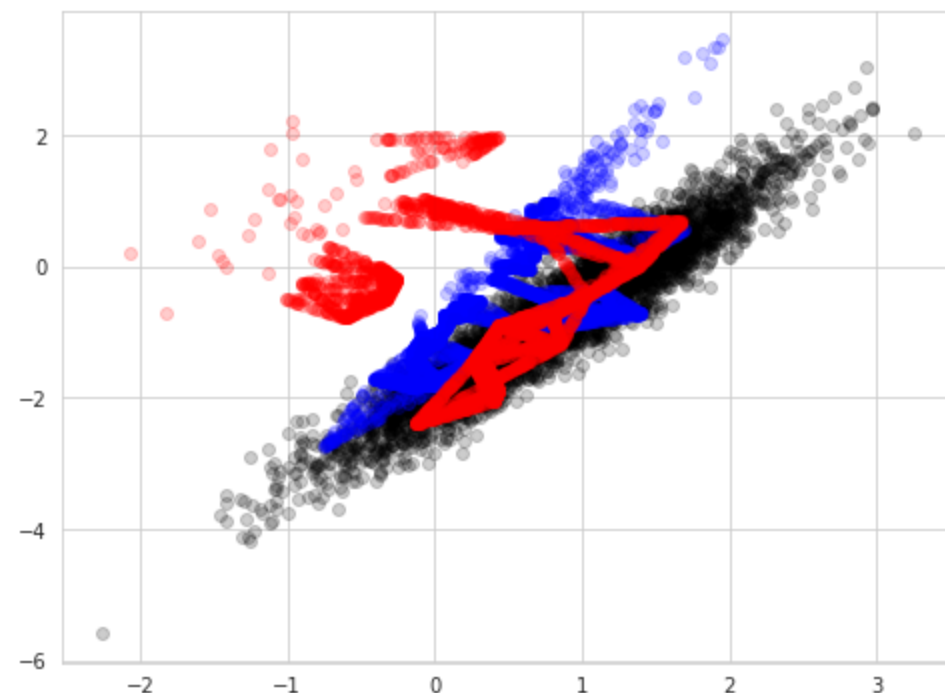
- Create new synthetic points between existing points *where classes overlap*

```
In [55]: from imblearn.over_sampling import ADASYN
```

```
X_adasyn, y_adasyn = ADASYN().fit_resample(X_imb, y_imb)  
Counter(y_adasyn).items()
```

```
Out[55]: dict_items([(2, 4674), (1, 4662), (0, 4673)])
```

```
In [56]: fig, ax = plt.subplots(1, 1, figsize=(8, 6))  
ax.scatter(X_adasyn[y_adasyn==2, 0], X_adasyn[y_adasyn==2, 1], c='k', alpha=.2);  
ax.scatter(X_adasyn[y_adasyn==1, 0], X_adasyn[y_adasyn==1, 1], c='b', alpha=.2);  
ax.scatter(X_adasyn[y_adasyn==0, 0], X_adasyn[y_adasyn==0, 1], c='r', alpha=.2);
```



Other methods for dealing with imbalanced classes

- Adjust class weight (sklearn)
 - Adjust decision threshold (sklearn)
 - Treat as anomaly detection
 - Generate/buy more labels
-
- See https://imbalanced-learn.readthedocs.io/en/stable/auto_examples/over-sampling/plot_comparison_over_sampling.html

Questions re Imbalanced Classes?

Data Processing and Delivery: ETL

- **Extract Transform Load**
- Extract: Reading in data
- Transform: Transforming data
- Load: Delivering data

Extract: Various Data Sources

- flatfiles (csv, excel)
 - semi-structured documents (json, html)
 - unstructured documents
 - data + schema (dataframe, database, parquet)
 - APIs (wikipedia, twitter, spotify, etc.)
 - databases
-
- Pandas to the rescue!
 - Plus other specialized libraries

Extracting Data with Pandas

- read_csv
- read_excel
- read_parquet
- read_json
- read_html
- read_sql
- read_clipboard
- ...

Extract Data: CSV

Comma Separated Values

Extract Data: CSV

Comma Separated Values

```
In [57]: %cat ../data/example.csv
```

```
Year,Make,Model,Description,Price
1997,Ford,E350,"ac, abs, moon",3000.00
1999,Chevy,"Venture Extended Edition","",4900.00
1999,Chevy,"Venture Extended Edition, Very Large",,5000.00
1996,Jeep,Grand Cherokee,"MUST SELL! air, moon roof, loaded",4799.00
```

Extract Data: CSV

Comma Separated Values

```
In [57]: %cat ../data/example.csv
```

```
Year,Make,Model,Description,Price
1997,Ford,E350,"ac, abs, moon",3000.00
1999,Chevy,"Venture Extended Edition","",4900.00
1999,Chevy,"Venture Extended Edition, Very Large",,5000.00
1996,Jeep,Grand Cherokee,"MUST SELL! air, moon roof, loaded",4799.00
```

```
In [58]: df = pd.read_csv('../data/example.csv',header=0,sep=',')
df.head()
```

Out[58]:

	Year	Make	Model	Description	Price
0	1997	Ford	E350	ac, abs, moon	3000.0
1	1999	Chevy	Venture Extended Edition	NaN	4900.0
2	1999	Chevy	Venture Extended Edition, Very Large	NaN	5000.0
3	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799.0

Extract Data: Excel

	A	B	C	D	E
1	Year	Make	Model	Description	Price
2	1997	Ford	E350	ac, abs, moon	3000
3	1999	Chevy	Venture Extended Edition		4900
4	1999	Chevy	Venture Extended Edition, Very Large		5000
5	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799

Extract Data: Excel

	A	B	C	D	E
1	Year	Make	Model	Description	Price
2	1997	Ford	E350	ac, abs, moon	3000
3	1999	Chevy	Venture Extended Edition		4900
4	1999	Chevy	Venture Extended Edition, Very Large		5000
5	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799

In [59]: `pd.read_excel('../data/example.xls')`

Out[59]:

	Year	Make	Model	Description	Price
0	1997	Ford	E350	ac, abs, moon	3000
1	1999	Chevy	Venture Extended Edition	NaN	4900
2	1999	Chevy	Venture Extended Edition, Very Large	NaN	5000
3	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799

Extract Data: Parquet

- open source column-oriented data storage
- part of the Apache Hadoop ecosystem
- often used when working with Spark
- requires additional parsing engine eg `pyarrow`
- includes both data and **schema**
- **Schema** : metadata about the dataset (column names, datatypes, etc.)

Extract Data: Parquet

- open source column-oriented data storage
- part of the Apache Hadoop ecosystem
- often used when working with Spark
- requires additional parsing engine eg `pyarrow`
- includes both data and **schema**
- **Schema** : metadata about the dataset (column names, datatypes, etc.)

```
In [60]: # conda install -n eods-s21 pyarrow
pd.read_parquet('../data/example.parquet')
```

Out[60]:

	Year	Make	Model	Description	Price
0	1997	Ford	E350	ac, abs, moon	3000.0
1	1999	Chevy	Venture Extended Edition	None	4900.0
2	1999	Chevy	Venture Extended Edition, Very Large	None	5000.0
3	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799.0

Extract Data: JSON

- JavaScript Object Notation
- often seen as return from api call
- looks like a dictionary or list of dictionaries
- pretty print using `json.loads(json_string)`

```
{  
  "firstName": "John",  
  "lastName": "Smith",  
  "isAlive": true,  
  "age": 27,  
  "phoneNumbers": [  
    {  
      "type": "home",  
      "number": "212 555-1234"  
    },  
    {  
      "type": "cell",  
      "number": "312 555-1234"  
    }  
  ]  
}
```

Extract Data: JSON

Extract Data: JSON

```
In [61]: json = """
{"0": {"Year": 1997,
      "Make": "Ford",
      "Model": "E350",
      "Description": "ac, abs, moon",
      "Price": 3000.0},
"1": {"Year": 1999,
      "Make": "Chevy",
      "Model": "Venture Extended Edition",
      "Description": null,
      "Price": 4900.0},
"2": {"Year": 1999,
      "Make": "Chevy",
      "Model": "Venture Extended Edition, Very Large",
      "Description": null,
      "Price": 5000.0},
"3": {"Year": 1996,
      "Make": "Jeep",
      "Model": "Grand Cherokee",
      "Description": "MUST SELL! air, moon roof, loaded",
      "Price": 4799.0}}
"""
```

Extract Data: JSON

```
In [61]: json = """
{"0": {"Year": 1997,
      "Make": "Ford",
      "Model": "E350",
      "Description": "ac, abs, moon",
      "Price": 3000.0},
"1": {"Year": 1999,
      "Make": "Chevy",
      "Model": "Venture Extended Edition",
      "Description": null,
      "Price": 4900.0},
"2": {"Year": 1999,
      "Make": "Chevy",
      "Model": "Venture Extended Edition, Very Large",
      "Description": null,
      "Price": 5000.0},
"3": {"Year": 1996,
      "Make": "Jeep",
      "Model": "Grand Cherokee",
      "Description": "MUST SELL! air, moon roof, loaded",
      "Price": 4799.0}}
"""
```

```
In [62]: pd.read_json(json,orient='index')
```

Out[62]:

	Year	Make	Model	Description	Price
0	1997	Ford	E350	ac, abs, moon	3000
1	1999	Chevy	Venture Extended Edition	None	4900
2	1999	Chevy	Venture Extended Edition, Very Large	None	5000
3	1996	Jeep	Grand Cherokee	MUST SELL! air, moon roof, loaded	4799

Extract Data: HTML

- HyperText Markup Language
- Parse with BeautifulSoup

Extract Data: HTML

- HyperText Markup Language
- Parse with BeautifulSoup

```
In [63]: html = """
<html>
  <head>
    <title>Example</title>
  </head>
  <body>
    <p id="first" class="example"><strong>Example text!</strong></p>
    <p id="second" class="example">And More!</p>
  </body>
</html>
"""

from bs4 import BeautifulSoup

soup = BeautifulSoup(html)
[p.text for p in soup('p')]
```

```
Out[63]: ['Example text!', 'And More!']
```

Extract Data: APIs

- Application **P**rogramming Interface
 - defines interactions between software components and resources
 - most datasources have an API
 - some require authentication
 - python libraries exist for most common APIs
-
- **requests**: library for making web requests and accessing the results

API Example: Wikipedia

API Example: Wikipedia

```
In [64]: import requests
url = 'http://en.wikipedia.org/w/api.php?action=query&prop=info&format=json&titles='
title = 'Data Science'
title = title.replace(' ', '%20')
print(url+title)
```

```
http://en.wikipedia.org/w/api.php?action=query&prop=info&format=json&titles=Data%20Science
```

API Example: Wikipedia

```
In [64]: import requests
url = 'http://en.wikipedia.org/w/api.php?action=query&prop=info&format=json&titles='
title = 'Data Science'
title = title.replace(' ', '%20')
print(url+title)
```

http://en.wikipedia.org/w/api.php?action=query&prop=info&format=json&titles=Data%20Science

```
In [65]: resp = requests.get(url+title)
resp.json()
```

```
Out[65]: {'batchcomplete': '',
'query': {'pages': {'49495124': {'pageid': 49495124,
'ns': 0,
'title': 'Data Science',
'contentmodel': 'wikitext',
'pagelanguage': 'en',
'pagelanguagehtmlcode': 'en',
'pagelanguagedir': 'ltr',
'touched': '2021-11-13T20:22:16Z',
'lastrevid': 706007296,
'length': 26,
'redirect': '',
'new': ''}}}}
```

```
In [66]: resp.text
```

```
Out[66]: '{"batchcomplete":"","query":{"pages":{"49495124":{"pageid":49495124,"ns":0,"title":"Data Science","contentmodel":"wikitext","pagelanguage":"en","pagelanguagehtmlcode":"en","pagelanguagedir":"ltr","touched":"2021-11-13T20:22:16Z","lastrevid":706007296,"length":26,"redirect":"","new":""}}}}'
```

API Example: Twitter

1. Apply for Twitter developer account
2. Create a Twitter application to generate tokens and secrets

API Example: Twitter

1. Apply for Twitter developer account
2. Create a Twitter application to generate tokens and secrets

```
In [67]: with open('/home/bgibson/proj/twitter/twitter_consumer_key.txt') as f:
          consumer_key = f.read().strip()
          with open('/home/bgibson/proj/twitter/twitter_consumer_secret.txt') as f:
              consumer_secret = f.read().strip()
          with open('/home/bgibson/proj/twitter/twitter_access_token.txt') as f:
              access_token = f.read().strip()
          with open('/home/bgibson/proj/twitter/twitter_access_token_secret.txt') as f:
              access_token_secret = f.read().strip()

          from twython import Twython
          twitter = Twython(consumer_key, consumer_secret, access_token, access_token_secret)
```

API Example: Twitter

1. Apply for Twitter developer account
2. Create a Twitter application to generate tokens and secrets

```
In [67]: with open('/home/bgibson/proj/twitter/twitter_consumer_key.txt') as f:
        consumer_key = f.read().strip()
        with open('/home/bgibson/proj/twitter/twitter_consumer_secret.txt') as f:
            consumer_secret = f.read().strip()
        with open('/home/bgibson/proj/twitter/twitter_access_token.txt') as f:
            access_token = f.read().strip()
        with open('/home/bgibson/proj/twitter/twitter_access_token_secret.txt') as f:
            access_token_secret = f.read().strip()

        from twython import Twython
        twitter = Twython(consumer_key, consumer_secret, access_token, access_token_secret)
```

```
In [68]: public_tweets = twitter.search(q='columbia')['statuses']
        for status in public_tweets[:3]:
            print('-----')
            print(status["text"])
```

```
-----
James Monroe is an American politician and actor who served as the most important invention of the original District of Columbi
a, the (1/2)
-----
And not Columbia, TN... ya small minded gyals ☹☹
-----
RT @ellisemmagarey: Workers at @Columbia are on their SIXTH week of strike & face mass retaliatory firings. This is their F
OURTH strike in...
```

API Example: Twitter

API Example: Twitter

```
In [69]: public_tweets[0]
```

```
Out[69]: {'created_at': 'Mon Dec 06 19:57:38 +0000 2021',
          'id': 1467946353493872643,
          'id_str': '1467946353493872643',
          'text': 'James Monroe is an American politician and actor who served as the most important invention of the original District of Columbia, the (1/2)',
          'truncated': False,
          'entities': {'hashtags': [], 'symbols': [], 'user_mentions': [], 'urls': []},
          'metadata': {'iso_language_code': 'en', 'result_type': 'recent'},
          'source': '<a href="http://www.github.com/ldermer/" rel="nofollow">dunning kruger bot</a>',
          'in_reply_to_status_id': None,
          'in_reply_to_status_id_str': None,
          'in_reply_to_user_id': None,
          'in_reply_to_user_id_str': None,
          'in_reply_to_screen_name': None,
          'user': {'id': 739988612243062784,
                   'id_str': '739988612243062784',
                   'name': 'dunningkrugerb主',
                   'screen_name': 'bottingkruger',
                   'location': 'Seattle, WA',
                   'description': "Ask me about a person, and I'll tell you everything I think I know. I only read Wikipedia, but I think I got this. Maintained by @lauriedermer.",
                   'url': None,
                   'entities': {'description': {'urls': []}},
                   'protected': False,
                   'followers_count': 71,
                   'friends_count': 1,
                   'listed_count': 8,
                   'created_at': 'Tue Jun 07 01:13:28 +0000 2016',
                   'favourites_count': 5,
                   'utc_offset': None,
                   'time_zone': None,
                   'geo_enabled': False,
                   'verified': False,
                   'statuses_count': 113407,
```

Transforming Data

- Standardization
 - Creating dummy variables
 - Filling missing data
 - One-Hot-Encoding
 - Binning
 - Parsing natural language
 - Dimensionality reduction
 - etc...
-
- Pipeline and ColumnTransformer

Extract and Transform Example: Titanic

Extract and Transform Example: Titanic

```
In [70]: titanic_url = ('https://raw.githubusercontent.com/amueller/'
                        'scipy-2017-sklearn/091d371/notebooks/datasets/titanic3.csv')
df_titanic = pd.read_csv(titanic_url)[['age', 'fare', 'embarked', 'sex', 'pclass', 'survived']]
display(df_titanic.head(1))
df_titanic.info()
```

	age	fare	embarked	sex	pclass	survived
0	29.0	211.3375	S	female	1	1

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1046 non-null   float64
1   fare        1308 non-null   float64
2   embarked    1307 non-null   object
3   sex         1309 non-null   object
4   pclass      1309 non-null   int64
5   survived    1309 non-null   int64
dtypes: float64(2), int64(2), object(2)
memory usage: 61.5+ KB
```

Extract and Transform Example: Titanic

```
In [70]: titanic_url = ('https://raw.githubusercontent.com/amueller/'
                      'scipy-2017-sklearn/091d371/notebooks/datasets/titanic3.csv')
df_titanic = pd.read_csv(titanic_url)[['age', 'fare', 'embarked', 'sex', 'pclass', 'survived']]
display(df_titanic.head(1))
df_titanic.info()
```

	age	fare	embarked	sex	pclass	survived
0	29.0	211.3375	S	female	1	1

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1309 entries, 0 to 1308
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1046 non-null   float64
1   fare        1308 non-null   float64
2   embarked    1307 non-null   object
3   sex         1309 non-null   object
4   pclass      1309 non-null   int64
5   survived    1309 non-null   int64
dtypes: float64(2), int64(2), object(2)
memory usage: 61.5+ KB
```

```
In [71]: X_titanic = df_titanic.drop('survived', axis=1)
y_titanic = df_titanic['survived']
X_titanic_train, X_titanic_test, y_titanic_train, y_titanic_test = train_test_split(X_titanic,
                                                                                      y_titanic,
                                                                                      test_size=0.2,
                                                                                      random_state=42)
```

Extract and Transform Example: Titanic

Extract and Transform Example: Titanic

```
In [72]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

numeric_features = ['age', 'fare']
numeric_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='median')),
                                      ('scaler', StandardScaler())])

categorical_features = ['embarked', 'sex', 'pclass']
categorical_transformer = Pipeline(steps=[('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                          ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(transformers=[('num', numeric_transformer, numeric_features),
                                             ('cat', categorical_transformer, categorical_features)])

pipe = Pipeline(steps=[('preprocessor', preprocessor),
                       ('classifier', LogisticRegression(solver='lbfgs', random_state=42))])

param_grid = {
    'preprocessor__num__imputer__strategy': ['mean', 'median'],
    'classifier__C': [0.1, 1.0, 10, 100],
}
gs_pipeline = GridSearchCV(pipe, param_grid, cv=3)
gs_pipeline.fit(X_titanic_train, y_titanic_train)
print("best test set score from grid search: {:.3f}".format(gs_pipeline.score(X_titanic_test, y_titanic_test)))
print("best parameter settings: {}".format(gs_pipeline.best_params_))
```

```
best test set score from grid search: 0.771
best parameter settings: {'classifier__C': 100, 'preprocessor__num__imputer__strategy': 'median'}
```

Loading Data with pandas

- to_csv
- to_excel
- to_json
- to_html
- to_parquet
- to_sql
- to_clipboard
- to_pickle

Delivering Data With Flask

Delivering Data With Flask

- Flask : lightweight web server
- can be used to create a small API to:
 - return transformed data
 - return predictions
 - return datasets
 - ...

Aside: Running python scripts from the command line

Aside: Running python scripts from the command line

In [73]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
# will run as script or on import
run_or_imported_at = datetime.now()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s}")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

Aside: Running python scripts from the command line

In [73]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
# will run as script or on import
run_or_imported_at = datetime.now()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s}")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

In [74]: `import sys`
`sys.path.append('../src/')`

`import sample_script`

```
this was run or imported at 2021-12-06 14:57:49.844899
__name__ = sample_script
```

Aside: Running python scripts from the command line

In [73]: `!cat ../src/sample_script.py`

```
# import necessary libraries and function
from datetime import datetime

# python as usual
# will run as script or on import
run_or_imported_at = datetime.now()
print(f"this was run or imported at {run_or_imported_at}")
print(f"__name__ = :s}")

if __name__ == "__main__":
    # will only run if this is a script
    # won't be run if imported
    print("running as a script")
```

In [74]: `import sys`
`sys.path.append('../src/')`

`import sample_script`

```
this was run or imported at 2021-12-06 14:57:49.844899
__name__ = sample_script
```

In [75]: `print(sample_script.run_or_imported_at)`

```
2021-12-06 14:57:49.844899
```

Aside: Function Decorators

- act like wrappers around functions
- decorators are prefixed by the "@" symbol
- placed above the function to be wrapped

Aside: Function Decorators

- act like wrappers around functions
- decorators are prefixed by the "@" symbol
- placed above the function to be wrapped

```
In [76]: def my_decorator(func):
          def wrapper():
              print("Happens before the function is called.")
              func()
              print("Happens after the function is called.")
          return wrapper

          @my_decorator
          def say_hello():
              print("Hello")

          say_hello()
```

```
Happens before the function is called.
Hello
Happens after the function is called.
```

Creating APIs: Flask

Creating APIs: Flask

```
In [77]: !cat ../src/hello_flask.py

from flask import Flask, escape, request

app = Flask(__name__)

@app.route('/')
def hello():
    name = request.args.get("name", "World")
    return f'Hello, {escape(name)}!\n'

if __name__ == '__main__':
    app.run()
```


Creating APIs: Flask

```
In [77]: !cat ../src/hello_flask.py

from flask import Flask, escape, request

app = Flask(__name__)

@app.route('/')
def hello():
    name = request.args.get("name", "World")
    return f'Hello, {escape(name)}!\n'

if __name__ == '__main__':
    app.run()
```

1. at command line, run: `$ python hello_flask.py`

2. in ipython (or notebook)

```
import requests
r = requests.get('http://127.0.0.1:5000/?name=Bryan')
print(r.text)
```

Creating APIs: Flask with Multiple Routes

Creating APIs: Flask with Multiple Routes

In [78]: !cat ../src/die_flask.py

```
import numpy as np
from flask import Flask, request, jsonify

app = Flask(__name__)

@app.route("/")
def help():
    return "Give the number of sides the die should have.\n"

@app.route("/<int:sides>")
def roll_die(sides):
    return str(np.random.randint(1,sides+1))

@app.route("/json/<int:sides>")
def roll_die_json(sides):
    return jsonify({'sides': sides, 'roll': np.random.randint(1,sides+1)})

if __name__ == '__main__':
    app.run()
```

GET vs POST

- **GET**: pass information in the url

```
127.0.0.1:5000/?firstname=Bryan&lastname=Gibson
```

- **POST**: pass information as additional http request (often JSON)

```
127.0.0.1:5000/  
{ 'firstname': 'Bryan', 'lastname': 'Gibson' }
```

Creating APIs: Flask

- Export trained models (and other data structures) using `pickle`

Creating APIs: Flask

- Export trained models (and other data structures) using `pickle`

```
In [79]: import pickle as pkl  
         with open('../data/titanic_pipeline_clf.pkl', 'wb') as f:  
             pkl.dump(gs_pipeline, f)
```

Creating APIs: Deliver Predictions Using Flask

Creating APIs: Deliver Predictions Using Flask

In [80]: !cat ../src/titanic_clf.py

```
from flask import Flask, escape, request, jsonify
import pickle as pkl
import pandas as pd

# need to train and pickle classifier first
with open('../data/titanic_pipeline_clf.pkl','rb') as f:
    clf = pkl.load(f)

app = Flask(__name__)

@app.route('/',methods=['POST'])
def predict():
    prediction = None
    query = pd.DataFrame(request.form,index=[0])
    print(query,flush=True)
    if query is not None:
        prediction = clf.predict(query)
    if prediction:
        return jsonify([str(x) for x in prediction])
    else:
        return 'no predictions made'

if __name__ == '__main__':
    app.run()
```


Creating APIs: Deliver Predictions Using Flask Cont.

Creating APIs: Deliver Predictions Using Flask Cont.

```
In [81]: query_label = df_titanic.iloc[0].loc['survived']
```

Creating APIs: Deliver Predictions Using Flask Cont.

```
In [81]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [82]: query = df_titanic.iloc[0, :-1].to_dict()  
query
```

```
Out[82]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

Creating APIs: Deliver Predictions Using Flask Cont.

```
In [81]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [82]: query = df_titanic.iloc[0, :-1].to_dict()  
query
```

```
Out[82]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

```
In [83]: query_label
```

```
Out[83]: 1
```

Creating APIs: Deliver Predictions Using Flask Cont.

```
In [81]: query_label = df_titanic.iloc[0].loc['survived']
```

```
In [82]: query = df_titanic.iloc[0, :-1].to_dict()  
query
```

```
Out[82]: {'age': 29.0, 'fare': 211.3375, 'embarked': 'S', 'sex': 'female', 'pclass': 1}
```

```
In [83]: query_label
```

```
Out[83]: 1
```

```
In [84]: # Start script from command line using: python titanic_clf.py  
# Then run the following in from the notebook  
#requests.post('http://127.0.0.1:5000/', data=query).text
```

Data Processing Summary

- ETL
- reading datafiles using pandas
- website scraping (requests,BeautifulSoup)
- accessing data via API
- Tranforming data with Pipelines
- Exposing data via API (Flask)

Questions re Data Processing and Delivery?