

Week_05_Quiz-qm2162

October 13, 2021

1 Week 5 Quiz

1.1 Qi Meng - qm2162

1.1.1 Due Sunday Oct 17th, 11:59pm

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('darkgrid')
%matplotlib inline
```

```
[2]: # Sklearn provides a set of commonly used example datasets.
# They can be accessed through the datasets submodule.
from sklearn import datasets

# We're going to use the Linnerud dataset to practice Regression in sklearn.

# The Linnerud dataset is a tiny multi-output regression dataset. It consists
# of three exercise (data) and three physiological (target) variables
# collected from twenty middle-aged men in a fitness club.
linnerud = datasets.load_linnerud()

# The features of the dataset contain data on 3 exercises
# Chins - number of chinups
# Situps - number of situps
# Jumps - number of jumping jacks

# Note that the features and target come as numpy matrices.
# We'll first load the features into a pandas dataframe.
df = pd.DataFrame(linnerud.data, columns=linnerud.feature_names)

# We'll also add the target to our dataframe.
# Note also that this dataset contains multiple targets.
# We'll only consider one of them: Weight
df['Weight'] = linnerud.target[:, linnerud.target_names.index('Weight')]
```

```
# For more information on the dataset, uncomment the print command below
#print(linnerud.DESCR)

# print the first 3 rows
df.head(3)
```


```
[2]:
```

	Chins	Situps	Jumps	Weight
0	5.0	162.0	60.0	191.0
1	2.0	110.0	60.0	189.0
2	12.0	101.0	101.0	193.0

```
[3]: # What is the relationship between Situps and Weight?

# Plot a scatterplot and best-fit line for x=Situps vs y=Weight
# using seaborn sns.regplot()
sns.regplot(x='Situps', y='Weight', data=df)
```

```
[3]: <AxesSubplot:xlabel='Situps', ylabel='Weight'>
```



```
[4]: # The above plot should indicate a negative relationship
# between Situps and Weight
# How much does Weight go down if Situps goes up?
# To answer this we'll train a simple linear model.
```

```

# First import LinearRegression from sklearn.linear_model
from sklearn.linear_model import LinearRegression

# Create a variable X containing the independent variable 'Situps'
# Note that sklearn expects X to be two dimensional
# so you must use one of the methods discussed in class
# to return a two dimensional object
X = df.Situps.values.reshape(-1, 1)

# Create a variable y containing the dependent variable 'Weight'
# Note that y should only be one dimensional,
# so a Series (single column of a dataframe) works fine here
y = df.Weight

# Instantiate a LinearRegression object with default parameter settings
# and store as lr
lr = LinearRegression()

# Fit lr using the X and y defined above
lr.fit(X=X, y=y)

# Using the learned parameters in coef_ and intercept_,
# by how much do we expect Weight to go down when Situps goes up by 1?
# Print with a precision of 2
print(f"We expect Weight to go down by {np.abs(lr.coef_[0]):0.2f} when Situps_
    ↳ goes up by 1.")

# Using the learned parameters in coef_ and intercept_,
# what should we expect weight to be when when Situps is 0?
# Print with a precision of 2
print(f"We should expect weight to be {lr.intercept_:0.2f} when Situps is 0.")

```

We expect Weight to go down by 0.19 when Situps goes up by 1.
 We should expect weight to be 206.92 when Situps is 0.

```

[5]: # How is Weight related to all 3 features?

# Create a list containing the 3 feature names we're interested in
# as strings: Chins, Situps, Jumps
# Store as feature_names
# We'll do this to make sure we don't include 'Weight' in the
# regression as an independent variable
feature_names = ['Chins', 'Situps', 'Jumps']

# Instantiate a second LinearRegression model with default parameters
# and store as mlr
# Fit this model using all of the columns in feature_names

```

```

# Note that here we can pass the 2D dataframe without needing to .reshape(),  

↳ etc.
mlr = LinearRegression()
mlr.fit(df[feature_names], y=y)

# For each feature name in feature_names, print out the name and  

#     corresponding learned coefficient
# It looks like one of the features actually has a positive relationship.
# Print coefficient values with a precision of 2.
for (name, coef) in zip(feature_names, mlr.coef_):
    print(f'{name:10s} : {coef:0.2f}')

```

```

Chins      : -0.48
Situps     : -0.22
Jumps      : 0.09

```

```

[6]: # NOT REQUIRED

# For those that are interested exploring how statsmodels works

# Import the statsmodels api as sm
import statsmodels.api as sm

# Store the 3 features from df as X
X = df[feature_names].copy()

# Add a constant to X (in order to learn the bias term) using sm.add_constant()
sm.add_constant(X)

# Instantiate and fit an OLS model using X and df.Weight as y
#     and store as sm_model
# Note that in OLS, the target y is the first parameter!
sm_model = sm.OLS(y, X).fit()

# Display the model summary
# Note that the coefficients in the summary match the values
#     found above using sklearn
sm_model.summary()

```

```

/Users/mengqi/opt/anaconda3/envs/eods-f21/lib/python3.8/site-
packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of
pandas all arguments of concat except for the argument 'objs' will be keyword-
only

```

```

    x = pd.concat(x[:, :order], 1)

```

```
[6]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
=====
Dep. Variable:                  Weight    R-squared (uncentered):
0.791
Model:                        OLS      Adj. R-squared (uncentered):
0.755
Method:                      Least Squares    F-statistic:
21.50
Date:                        Wed, 13 Oct 2021    Prob (F-statistic):
5.07e-06
Time:                        03:40:45    Log-Likelihood:
-116.59
No. Observations:            20    AIC:
239.2
Df Residuals:                17    BIC:
242.2
Df Model:                    3
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
Chins                1.6422     5.376     0.305     0.764    -9.701    12.985
Situps               0.9735     0.442     2.201     0.042     0.041     1.906
Jumps              -0.1295     0.535    -0.242     0.812    -1.259     1.000
=====
Omnibus:                0.243    Durbin-Watson:           1.462
Prob(Omnibus):          0.886    Jarque-Bera (JB):         0.412
Skew:                   0.185    Prob(JB):                 0.814
Kurtosis:               2.402    Cond. No.                 47.8
=====

Notes:
[1] R2 is computed without centering (uncentered) since the model does not
contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```