

Week_06_Quiz-qm2162

October 24, 2021

1 Week 6 Quiz

1.1 Qi Meng - qm2162

1.1.1 Due Sunday Oct 24 11:59pm ET

In this quiz we'll be loading some data, training a few models and plotting their decision boundaries to compare how the models perform.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn import datasets
from mlxtend.plotting import plot_decision_regions

sns.set_style('darkgrid')
%matplotlib inline

[2]: # For this quiz we'll be using the "iris" dataset, a very common dataset seen
# when learning about classification.

# Using sklearn.datasets, load the dataset into iris
iris = datasets.load_iris()

[3]: # As we've seen in other datasets loaded from sklearn, the data is contained in
# iris.data and feature names in iris.feature_names.
# Create a new pd.DataFrame X_iris that contains iris.data
# with columns=iris.feature_names
X_iris = pd.DataFrame(iris.data, columns=iris.feature_names)

# In order to make the data easier to plot, keep only the first 2 features:
# "sepal length (cm)", "sepal width (cm)"
# Store back into X_iris
X_iris = X_iris.iloc[:, :2]

# We'll check to make sure that X_iris only has 2 columns
assert X_iris.shape == (150, 2)
```

```

# The target/label is contained in iris.target.
# Create a pd.Series y_iris containing iris.target.
y_iris = pd.Series(iris.target)

# Print out the number of rows per label.
# Note that this is a multiclass problem (3 classes)
#     each with 50 observations
y_iris.value_counts()

```

```

[3]: 0    50
      1    50
      2    50
      dtype: int64

```

```

[4]: # Import LogisticRegression from sklearn.linear_model
      from sklearn.linear_model import LogisticRegression

      # Instantiate the LogisticRegression model with default settings
      # and fit on X_iris and y_iris
      # Store in logreg
      logreg = LogisticRegression()
      logreg.fit(X_iris, y_iris)


      # Plot the training set and trained classifier with
      # plot_decision_regions() which takes numpy arrays
      # so use X_iris.values, y_iris.values
      plot_decision_regions(X_iris.values, y_iris.values, logreg)

```

```

[4]: <AxesSubplot:>

```




```
[5]: # Import KNeighborsClassifier from sklearn.neighbors
from sklearn.neighbors import KNeighborsClassifier

# Instantiate the KNeighborsClassifier model with default settings
# and fit on X_iris and y_iris
# Store in knn
knn = KNeighborsClassifier()
knn.fit(X_iris, y_iris)

# Plot the training set and trained classifier with plot_decision_regions()
plot_decision_regions(X_iris.values, y_iris.values, knn)
```

```
[5]: <AxesSubplot:>
```



```
[6]: # Import GradientBoostingClassifier from sklearn.ensemble
from sklearn.ensemble import GradientBoostingClassifier

# Instantiate the GradientBoostingClassifier with default settings
# and fit on X_iris and y_iris
# Store in gbc
gbc = GradientBoostingClassifier()
gbc.fit(X_iris, y_iris)

# Plot the training set and trained classifier with plot_decision_regions()
plot_decision_regions(X_iris.values, y_iris.values, gbc)
```

[6]: <AxesSubplot:>

